

Chapter 2

Quantum Hardware: High-Temperature Asymmetric-SQUID Qubit

Summary

This chapter presents a comprehensive theoretical analysis of asymmetric SQUID-based superconducting qubits across two operational paradigms. Part A introduces the architecture and theoretical framework. Part B validates baseline performance at 10-20 mK with 99.55% CNOT fidelity. Part C explores elevated-temperature operation at 300 GHz and 4-10 K, reducing cooling costs by 70-80% while maintaining 97-99% gate fidelities.

Part A

Asymmetric-SQUID Qubit

2.1 Status of the Art

Superconducting qubits have emerged as one of the most promising platforms for scalable quantum computing due to their lithographic compatibility, strong nonlinearity, and relatively high coherence times. The relaxation time T_1 denotes the characteristic time for a qubit to decay from the excited state $|1\rangle$ to the ground state $|0\rangle$ due to energy loss to its environment. The coherence time T_2 quantifies how long a qubit can preserve quantum phase information before decohering as a result of environmental fluctuations. The transmon qubit, introduced by Koch *et al.* [72], mitigates charge noise sensitivity by increasing the ratio E_J/E_C ¹, and has become the de facto standard in large-scale processors [3]. Recent devices have reached coherence times exceeding 0.3 ms [73] and, very lately, even beyond [11], though typically require operation below 20 mK in dilution refrigerators.

Alternative designs aim to enhance robustness or integrate new functionalities. Flux

¹The ratio $\frac{E_J}{E_C}$, where $E_J = \frac{\Phi_0 I_c}{2\pi}$, $E_C = \frac{e^2}{2C}$, is the figure of merit distinguishing different regimes of Josephson qubits. Here E_J is the Josephson energy determined by the junction critical current I_c , and E_C is the charging energy associated with the junction capacitance C .

qubits [74] offer fast tunability via magnetic flux, while capacitively shunted flux qubits (C-shunt) and fluxonium [75] achieve improved noise resilience. However, these designs still rely on ultra-low-temperature operation to suppress thermal excitations ($k_B T \ll \hbar \omega_q$), imposing significant cryogenic overheads.

High- T_c superconductors, such as NbN, have been explored for kinetic inductance detectors [76] and Josephson Junctions [77], showing potential for operation at above 4 K. Recent developments in epitaxial NbN/AlN/NbN junctions have demonstrated significant progress: Nakamura et al. [78] achieved $T_1 = 16.3 \mu\text{s}$ and $T_2, \text{echo} = 21.5 \mu\text{s}$ in all-nitride qubits on silicon substrates, with $T_c \approx 16$ K and superconducting gap $2\Delta \approx 5.2$ meV. While these coherence times of NbN-based Josephson Junction qubits remain below those of conventional Al/AlO_x/Al junctions ($\approx 300 \mu\text{s}$), as demonstrated by Nakamura et al. [79], the larger superconducting gap significantly suppresses quasiparticle generation. Such qubits could co-locate with cryogenic CMOS electronics [80], reducing control latency and wiring complexity.

Cavity quantum electrodynamics (cQED) provides a robust framework for qubit read-out and interaction [81]. Dispersive coupling to superconducting resonators enables high-fidelity, quantum non-demolition measurements [81, 82]. Multi-qubit gates, such as CNOTs, have been demonstrated with fidelities above 99% using microwave-activated interactions [12].

Experimental efforts toward elevated-temperature Josephson Junctions remain limited. Yamashita *et al.* demonstrated NbN-based Josephson devices operating up to 10 K [83]. Recent breakthroughs have significantly pushed the operating temperature frontier. Anferov et al. [84] demonstrated niobium-based transmon qubits operating at frequencies up to 24 GHz with coherence times of $\sim 1 \mu\text{s}$ at temperatures exceeding 200 mK, showing that quasiparticle decoherence can be suppressed up to approximately 1 K when using high-gap superconductors. This work aims to validate the fundamental approach of combining elevated frequencies with high- T_c materials for higher-temperature operation, though our target of 4-10 K at 300 GHz represents a further $\sim 40\times$ increase in operating temperature beyond the current state-of-the-art. More recent approaches propose hybrid systems, embedding High- T_c junctions in low-loss resonator environments, or leveraging materials such as NbTiN and granular aluminum for reduced surface loss [85]. The main obstacles remain dielectric losses, quasiparticle generation at higher temperatures, and maintaining anharmonicity for gate operations [72, 86, 87, 88, 89, 90].

In conclusion, while superconducting qubits at millikelvin temperatures have achieved state-of-the-art performance, their cryogenic requirements limit scalability and deployment. High- T_c approaches, though nascent, could unlock simpler cryogenic architectures and integration with cryo-CMOS, potentially enabling more compact and cost-effective quantum processors that could benefit both quantum computation [91] and quantum sensing [92] for HEP and Theoretical physics. Moreover, by exploiting the strong anharmonicity of an asymmetric-SQUID transmon [93, 94], the $|0\rangle$, $|1\rangle$, and $|2\rangle$ levels can be spectrally resolved, allowing the device to operate as a qutrit (qudit). Such multilevel control enables its use for simulating three-flavor neutrino oscillations, where each flavor state can be mapped onto a distinct qutrit level [95, 96, 97]. Further progress in the future will obviously depend on material engineering, specific assembly setup, and the development of high-frequency readout schemes compatible with elevated operating temperatures. These constraints motivate exploring alternative operating regimes, provided that coherence and controllability can be preserved.

2.2 The NbN-based elevated temperature transmon qubit project

The system under study is based on a capacitively coupled transmon-like qubit, implemented as an asymmetric dc SQUID to enable in-situ frequency tunability via magnetic flux bias. The asymmetry prevents the Josephson energy from vanishing exactly at half a flux quantum, thereby stabilizing the qubit frequency against flux noise at the so-called *sweet spot* [72, 79]. Hutchings et al. [78] demonstrated that weakly tunable asymmetric qubits can achieve flux-independent dephasing times (T_2 , echo $\approx 40\text{-}45\ \mu\text{s}$) over a ~ 340 MHz tunable range, validating this design approach.

For a SQUID with junction energies E_{J1} and E_{J2} , the effective flux-dependent Josephson energy is

$$E_J(\Phi) = \sqrt{E_\Sigma^2 \cos^2\left(\pi \frac{\Phi}{\Phi_0}\right) + E_\Delta^2 \sin^2\left(\pi \frac{\Phi}{\Phi_0}\right)}, \quad (2.1)$$

where $E_\Sigma = E_{J1} + E_{J2}$, $E_\Delta = E_{J1} - E_{J2}$, and Φ_0 is the flux quantum. This reduces to the symmetric SQUID result for $E_{J1} = E_{J2}$ ($E_\Delta = 0$), while for finite asymmetry it ensures a non-zero minimum Josephson energy at $\Phi = \Phi_0/2$. In the Symmetric case $E_\Delta = 0$, the usual SQUID modulation is recovered.

The Hamiltonian of qubit i is then of the transmon² type [72]:

$$H_{q,i} = 4E_{C,i}(\hat{n}_i - n_{g,i})^2 - E_{J,i}(\Phi_i) \cos \hat{\varphi}_i, \quad (2.2)$$

where $\hat{n}_i$ is the Cooper-pair number operator, $n_{g,i}$ the offset charge, and $\hat{\varphi}_i$ the superconducting phase across the junction.

2.2.1 Anharmonicity: Definition and Physical Origin

Starting from the Josephson potential energy

$$U_J(\phi) = -E_J \cos \phi, \quad (2.3)$$

and assuming small phase fluctuations around the potential minimum ($\langle \phi^2 \rangle \ll 1$), as justified in the transmon regime $E_J/E_C \gg 1$, the cosine potential can be expanded in a Taylor series around $\phi = 0$ up to fourth order:

$$-E_J \cos \phi \simeq \frac{E_J}{2} \phi^2 - \frac{E_J}{24} \phi^4 + \mathcal{O}(\phi^6). \quad (2.4)$$

The quadratic term corresponds to an effective harmonic oscillator, while the quartic correction provides the leading-order anharmonicity responsible for the non-equidistant energy level spacing of the transmon qubit.

The φ^2 term gives rise to a harmonic oscillator potential, while the φ^4 term introduces *negative* anharmonicity. After a canonical transformation to ladder operators ($a, a^\dagger$), the Hamiltonian becomes:

$$H \approx \hbar\omega_{01} \left(a^\dagger a + \frac{1}{2} \right) + \frac{\alpha}{2} (a^\dagger a) (a^\dagger a - 1) \quad (2.5)$$

²A qubit is said to operate in the *transmon* regime when the ratio of Josephson to charging energy is large, $E_J/E_C \gtrsim 50$. In this limit, the sensitivity of the qubit transition frequency to offset charge n_g is exponentially suppressed, while a finite anharmonicity $\alpha \approx -E_C$ remains to isolate the computational subspace. This defines the transmon character of the device.

where

$$\omega_{01} \approx \sqrt{8E_J E_C} - E_C \quad (2.6)$$

$$\alpha \approx -E_C \quad (2.7)$$

The **negative sign** in Eq. (2.7) arises from the fact that the Josephson cosine potential has a finite barrier height, unlike a pure harmonic potential which extends to infinity. Consequently, energy levels become progressively closer as one climbs the energy ladder:

$$\omega_{01} > \omega_{12} > \omega_{23} > \omega_{34} > \dots \quad (2.8)$$

This level bunching is characteristic of a Morse-like potential and is fundamental to transmon operation.

The negative anharmonicity reflects the “softening” of the potential at large phase excursions. As the system is excited to higher energy levels, the wavefunction extends into regions where the cosine potential deviates significantly from its harmonic (parabolic) approximation near the minimum. The reduced curvature at larger φ leads to decreased level spacings, manifesting as $\alpha < 0$.

2.2.2 Qubit characterization

Qubit characterization will be carried out at the theoretical and numerical level. The relevant metrics are the energy relaxation time T_1 and the dephasing time T_2 , related by

$$\frac{1}{T_2} = \frac{1}{2T_1} + \Gamma_\varphi, \quad (2.9)$$

where Γ_φ is the pure dephasing rate. For asymmetric SQUID-based qubits, modeling at the flux sweet spot predicts a suppression of flux-noise contributions to Γ_φ , thus enhancing coherence. Numerical simulations will be employed to estimate these coherence times and to identify optimal operating regimes. This approach allows us to benchmark expected performance and define feasibility criteria for subsequent experimental validation, following protocols such as Ramsey fringes, echo sequences, and Purcell-filtered spectroscopy as described in literature [98, 99].

2.2.3 Models and Simulations

To capture the dynamics of our coupled qubit–resonator system we employ the Lindblad master equation formalism, which provides a rigorous framework for open quantum systems interacting with their environment [81, 100]. This approach is essential since superconducting qubits are never perfectly isolated: energy relaxation, pure dephasing, and cavity photon loss all contribute to non-unitary dynamics.

The time evolution of the system density matrix ρ is governed by

$$\dot{\rho} = -\frac{i}{\hbar}[H, \rho] + \sum_k \mathcal{D}[c_k]\rho, \quad (2.10)$$

where $\mathcal{D}[c_k]\rho = c_k \rho c_k^\dagger - \frac{1}{2}\{c_k^\dagger c_k, \rho\}$ is the Lindblad dissipator. The operators c_k are the

collapse operators associated with physical decoherence channels:

$$c_{i,\downarrow} = \sqrt{\gamma_{1,i}(1 + \bar{n}_i)} \sigma_-^{(i)} \quad (\text{qubit relaxation}), \quad (2.11)$$

$$c_{i,\uparrow} = \sqrt{\gamma_{1,i}\bar{n}_i} \sigma_+^{(i)} \quad (\text{thermal excitation}), \quad (2.12)$$

$$c_{i,\phi} = \sqrt{\gamma_{\phi,i}} \sigma_z^{(i)} \quad (\text{pure dephasing}), \quad (2.13)$$

$$c_r = \sqrt{\kappa} a \quad (\text{resonator photon loss}). \quad (2.14)$$

Here $\gamma_{1,i} = 1/T_{1,i}$ and $\gamma_{\phi,i}$ are the qubit relaxation and pure dephasing rates, κ is the resonator linewidth, and $\bar{n}_i = (e^{\hbar\omega_{01,i}/k_B T} - 1)^{-1}$ is the thermal photon number at temperature T , for the i -th qubit.

Numerical integration of the master equation is carried out with the QuTiP library [101, 102], enabling both time-domain gate simulations (with DRAG pulse shaping³) and steady-state readout analysis. This allows us to extract coherence metrics, gate fidelities, and readout signal-to-noise ratios under realistic conditions.

Parameter ranges are chosen consistently with established experimental benchmarks for superconducting qubits in circuit QED: $\omega_{01}/2\pi \in [5, 8]$ GHz, anharmonicity $\alpha/2\pi \in [-0.10, -0.30]$ GHz, qubit-resonator coupling $g/2\pi \in [50, 150]$ MHz, resonator linewidth $\kappa/2\pi \in [0.2, 2]$ MHz, and operation temperature $T \in [4, 10]$ K. These values are compatible with transmon designs in the literature [2, 12, 98] and ensure operation in the dispersive regime with $\chi/\kappa \gg 1$. They also guarantee that the Purcell-limited relaxation time T_1 exceeds several microseconds, which is necessary for gate fidelities above the surface-code threshold.

The Lindblad framework, combined with QuTiP [102] simulations, will provide a quantitatively reliable method to predict the coherence properties, gate performance, and readout fidelity of our elevated-temperature NbN qubits.

2.2.4 Resonator coupling

Each transmon qubit is coupled with a planar superconducting resonator bus. The superconducting resonator bus of frequency ω_r is modeled as a harmonic oscillator:

$$H_r = \hbar\omega_r a^\dagger a. \quad (2.15)$$

Coupling between qubit i and the resonator is of dipole type:

$$H_{qr,i} = \hbar g_i (a + a^\dagger) \hat{n}_i, \quad (2.16)$$

In a multi-qubit system possible direct capacitive couplings between neighboring qubits give

$$H_{qq} = \sum_{\langle i,j \rangle} J_{ij} \hat{n}_i \hat{n}_j. \quad (2.17)$$

The total Hamiltonian for the multi-qubit system, in line with standard circuit-QED formulations [104, 105], is

$$H = H_r + \sum_{i=1}^N H_{q,i} + \sum_{i=1}^N H_{qr,i} + H_{qq}, \quad (2.18)$$

³The DRAG (Derivative Removal by Adiabatic Gate) technique is a pulse-shaping method designed to suppress leakage out of the computational subspace in weakly anharmonic qubits such as the transmon. By adding a corrective quadrature component proportional to the derivative of the main drive envelope, DRAG reduces unwanted population of higher levels (e.g. $|2\rangle$) and mitigates phase errors, thereby enhancing single-qubit gate fidelities [5, 103].

where the terms respectively describe the resonator, the qubits, their coupling to the resonator, and possible direct qubit–qubit interactions.

2.2.5 Dispersive regime

Conventions. Throughout these sections, we report coupling rates and dispersive shifts in *frequency units* (Hz), i.e., $g/2\pi$, $\Delta/2\pi$, $\alpha/2\pi$, $\chi/2\pi$, $\kappa/2\pi$. This avoids 2π ambiguities when comparing text, tables, and figure captions.

In the dispersive limit, defined by $|\Delta_i| = |\omega_{q,i} - \omega_r| \gg g_i$, the Hamiltonian can be diagonalized perturbatively [81]:

$$H_{\text{disp}} = \hbar \left(\omega_r + \sum_i \chi_i \sigma_z^{(i)} \right) a^\dagger a + \frac{\hbar}{2} \sum_i \tilde{\omega}_{q,i} \sigma_z^{(i)} + \sum_{\langle i,j \rangle} \zeta_{ij} \sigma_z^{(i)} \sigma_z^{(j)}, \quad (2.19)$$

where $\chi_i \simeq g_i^2 \alpha_i / [\Delta_i (\Delta_i + \alpha_i)]$ are the dispersive shifts (with anharmonicity α_i), and ζ_{ij} are effective cross-Kerr interactions.

At elevated temperature (4–10 K), the thermal occupation of the qubit transition

$$\bar{n}(T) = \frac{1}{\exp(\hbar\omega_{01}/k_B T) - 1} \quad (2.20)$$

must remain small. Our design must ensure $\hbar\omega_{01} \gg k_B T$, so that $\bar{n}(T) < 0.05$ even at 4 K.

2.2.6 Control electronics critical aspects and possible solutions

To keep thermal noise acceptable, we are forced to operate our transmons at frequencies of around 250 GHz to 300 GHz. A critical challenge in the present project is the generation of coherent control pulses at the qubit idle frequency, expected to lie in the 250–300 GHz range to ensure thermalization at 4 K. Standard quantum control electronics are designed for the 4–20 GHz regime, and extending their performance into the sub-THz domain is non-trivial. We identify three possible strategies:

- **Flux-dip control.** The qubit frequency $\omega_{01}(\Phi)$ can be dynamically shifted down with fast flux pulses during gate operations[106, 107, 108, 109, 110], lowering ω_{01} into the accessible 20–40 GHz range.
- **Commercial sub-THz signal sources.**
- **Frequency mixing and multiplication chains.** Employ frequency mixing and harmonic multiplication starting from a low-phase-noise local oscillator in the 10–20 GHz range.

The choice between flux-dip control, direct sub-THz sources, and frequency mixing will depend on the balance between cost, complexity, and simulation results.

A feasibility study will provide the optimal choice for the experimental phase of the project.

2.3 Comprehensive analysis of asymmetric-squid-based qubit

Part A has presented a comprehensive analysis of asymmetric SQUID-based superconducting qubits in two operational regimes. **Part B** demonstrates baseline performance at standard cryogenic temperatures (10–20 mK) with qubit frequencies of 5–6 GHz, achieving CNOT gate fidelities approaching 99% (98.7%). **Part C** explores a novel and ambitious extension: operating at elevated temperatures (4–10 K) by pushing qubit frequencies to 300 GHz. While technically challenging and unprecedented, this approach could revolutionize quantum computing by dramatically reducing cooling costs and enabling scalable deployment. We provide a theoretical framework, feasibility assessment, and a roadmap for achieving this breakthrough. Lastly, in **Part D**, we present the conclusions and some recommendations.

1. **Baseline validation:** Demonstrate that asymmetric SQUID qubits achieve state-of-the-art performance at conventional operating conditions
2. **Transformative innovation:** Explore the frontier of elevated-temperature quantum computing by pushing to 300 GHz operation at 4–10 K

The baseline provides proof of concept and validates our design approach. The innovation represents a high-risk, high-reward pathway that, if successful, would fundamentally change the landscape of quantum computing deployment.

Part B

Standard Cryogenic Operation

2.4 System Design and Parameters

We analyze a representative high- T_c NbN transmon with:

- Charging energy: $E_C = 0.2$ GHz
- Junction energies: $E_{J1} = 15.0$ GHz, $E_{J2} = 12.0$ GHz
- Asymmetry: $\eta = (E_{J1} - E_{J2})/(E_{J1} + E_{J2}) = 11.1\%$
- Ratio: $E_J/E_C|_{\Phi=0} = 135$ (firmly in transmon regime)

Spectrum at the Flux Sweet Spot

At $\Phi = \Phi_0/2$, the effective Josephson energy is $E_J(\Phi_0/2) = 3.000$ GHz. Numerical diagonalization with QuTiP [102] of the Hamiltonian (2.2) in the charge basis (using 30 charge states) yields:

Level	Energy (GHz)	Transition (GHz)
E_0	-1.001	–
E_1	+0.971	$\omega_{01} = 1.972$
E_2	+2.581	$\omega_{12} = 1.610$
E_3	+3.908	$\omega_{23} = 1.327$
E_4	+4.974	$\omega_{34} = 1.066$

Table 2.1: Energy eigenvalues and transition frequencies at the sweet spot.

The anharmonicity is:

$$\alpha = \omega_{12} - \omega_{01} = 1.610 - 1.972 = -0.362 \text{ GHz} = -362 \text{ MHz} \quad (2.21)$$

Comparing with the theoretical prediction $\alpha \approx -E_C = -200$ MHz, we find excellent qualitative agreement. The discrepancy arises because the approximate formula (2.7) is strictly valid only for $E_J/E_C \gg 1$, whereas at the sweet spot we have $E_J/E_C = 15$, which is moderate.




Figure 2.1: Asymmetric Squid energy vs flux.




Figure 2.2: Asymmetric Squid anharmonicity.

Anharmonicity as Percentage of Qubit Frequency

$$\frac{|\alpha|}{\omega_{01}} = \frac{362}{1972} \approx 18.4\% \quad (2.22)$$

This is within the acceptable range for transmon operation. We note that the relatively large $|\alpha|/\omega_{01}$ at the sweet spot enhances spectral selectivity, but it also indicates that the device is less deep in the transmon limit at that bias point, implying a trade-off between tunability, noise protection, and level structure. While typical Al-based transmons operate with $|\alpha|/\omega_{01} \sim 1 - 5\%$, the larger relative anharmonicity here is a consequence of the lower qubit frequency necessitated at this temperature regime of operation.

2.4.1 Significance of Negative Anharmonicity for Quantum Gates

Selective Addressing

Negative anharmonicity enables frequency-selective control. A microwave drive at frequency ω_{01} can excite the $|0\rangle \leftrightarrow |1\rangle$ transition without significantly affecting the $|1\rangle \leftrightarrow |2\rangle$ transition at frequency ω_{12} , provided:

$$\Omega_{\text{Rabi}} \ll |\alpha| \ll \omega_{01} \quad (2.23)$$

where Ω_{Rabi} is the gate Rabi frequency.

For our design:

$$\Omega_{\text{Rabi}} \lesssim \frac{|\alpha|}{3} \approx 120 \text{ MHz} \quad (2.24)$$

This corresponds to single-qubit gate times:

$$\tau_{\text{gate}} \sim \frac{\pi}{2\Omega_{\text{Rabi}}} \approx 13 \text{ ns} \quad (2.25)$$

Leakage Suppression

During gate operations, unwanted transitions to $|2\rangle$ (leakage) must be suppressed. The anharmonicity provides a natural energy detuning:

$$\Delta_{\text{leak}} = |\alpha| = 362 \text{ MHz} \quad (2.26)$$

For a gate pulse with bandwidth $\Delta\omega \sim 1/\tau_{\text{gate}} \approx 80 \text{ MHz}$, the leakage probability scales as:

$$P_{\text{leak}} \sim \left(\frac{\Delta\omega}{\Delta_{\text{leak}}} \right)^2 \approx \left(\frac{80}{362} \right)^2 \approx 5\% \quad (2.27)$$

This can be further reduced to $< 0.1\%$ using DRAG (Derivative Removal by Adiabatic Gate) pulse shaping [5], which explicitly compensates for the finite anharmonicity.

Two-Qubit Gate Feasibility

For two-qubit gates mediated by the cavity bus, two distinct interactions emerge from the Schrieffer–Wolff reduction. The dominant transverse exchange coupling (Majer *et al.* [?]) takes the form

$$J = \frac{g_1 g_2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right), \quad (2.28)$$

where g_i is the qubit–resonator coupling and $\Delta_i = \omega_{q,i} - \omega_r$ the qubit–cavity detuning. The longitudinal cross-Kerr coupling requires the transmon anharmonicity α and reads [81]

$$\zeta_{zz} = 2J^2 \left[\frac{1}{\Delta_q - \alpha_2} - \frac{1}{\Delta_q + \alpha_1} \right], \quad (2.29)$$

with $\Delta_q = \omega_{q,1} - \omega_{q,2}$ the qubit–qubit detuning. The two interactions have different physical origins and very different magnitudes: in the relevant operating regime J dominates over ζ_{zz} by an order of magnitude, with direct consequences for the choice of two-qubit gate protocol. The full derivation of both couplings, with numerical validation against exact diagonalization, is given in Section 2.10, while the gate-protocol analysis is presented in Section 2.11.

2.4.2 Thermal Population Analysis

Operating at standard temperature ($T = 10 \text{ mK}$ to 20 mK) ensures negligible thermal population of the excited state. The thermal occupation is:

$$\bar{n}(T) = \frac{1}{\exp(\hbar\omega_{01}/k_B T) - 1} \quad (2.30)$$

For $\omega_{01} = 1.972 \text{ GHz}$ and $k_B/\hbar = 20.836 \text{ GHz/K}$:

For comparison, state-of-the-art Al-based transmons typically operate at $T \sim 20 \text{ mK}$ where $\bar{n} < 10^{-6}$ for $\omega_{01} \sim 5 \text{ GHz}$.

To fully characterize the elevated-temperature operating regime, we analyze the thermal population landscape across a wide parameter space. Figure 2.3 presents a comprehensive six-panel analysis exploring the relationships between qubit frequency, temperature, and thermal excitation. The analysis reveals several key insights:




Figure 2.3: Comprehensive thermal population analysis for elevated-temperature operation. **(a)** Thermal population versus temperature for different qubit frequencies, showing the HEATS-Q operating range (10 mK to 20 mK shaded). **(b)** Thermal population versus qubit frequency for different temperatures, with the HEATS-Q sweet spot marked at 1.972 GHz. **(c)** Two-dimensional thermal population landscape showing contours of constant thermal excitation; the white dashed box indicates the HEATS-Q operating region. **(d)** Energy ratio $\hbar\omega_{01}/(k_B T)$ versus temperature, with thresholds for good (> 10), marginal (> 5), and poor (> 1) operation. **(e)** Performance comparison table showing thermal populations and energy ratios for traditional transmons versus HEATS-Q at various temperatures. **(f)** Required minimum qubit frequency to maintain specified thermal population thresholds as a function of operating temperature, with HEATS-Q design points overlaid.

Temperature (mK)	$\bar{n}(T)$
10	7.76×10^{-5}
12	3.76×10^{-4}
14	1.16×10^{-3}
16	2.70×10^{-3}
20	5.23×10^{-3}

Table 2.2: Thermal population at various temperatures.

Temperature Dependence. Panel (a) demonstrates that for a fixed qubit frequency, thermal population increases exponentially with temperature. The HEATS-Q design operating at $\omega_{01} = 1.972$ GHz maintains thermal excitation below 0.01% across the entire 4–10 K range, well within the acceptable threshold for high-fidelity quantum operations.

Frequency Requirements. Panel (b) shows that higher qubit frequencies dramatically reduce thermal populations at a given temperature. However, as discussed in **Part C**, frequencies above 40 GHz present significant control challenges. The HEATS-Q sweet-spot frequency of 1.972 GHz represents an optimal balance between thermal suppression and experimental accessibility.

Design Space Mapping. The contour plot in panel (c) provides a global view of the thermal landscape. The logarithmic color scale reveals that the HEATS-Q operating region (white dashed rectangle) lies comfortably within a low-thermal-population zone (blue region, $\bar{n} < 0.01$). This provides significant design margin and tolerance to parameter variations.

Energy Ratio Criterion. Panel (d) plots the dimensionless energy ratio $\hbar\omega_{01}/(k_B T)$, which must be $\gg 1$ for quantum behavior to dominate over thermal fluctuations. At the HEATS-Q sweet spot, this ratio exceeds nine even at 10 mK, satisfying the criterion for reliable qubit operation. For comparison, traditional transmons at 20 mK achieve ratios of $\sim 10^4$, but this extreme suppression is unnecessary for fault-tolerant quantum computing with modern error correction codes.

Performance Benchmarking. The comparison table in panel (e) quantitatively demonstrates that HEATS-Q thermal populations remain orders of magnitude below typical gate error rates (0.1–1%). While traditional Al-based transmons achieve negligible thermal excitation at 20 mK, the infrastructural cost of dilution refrigeration makes scaled deployment challenging. According to our estimation, the HEATS-Q approach sacrifices four orders of magnitude in thermal suppression but gains a 70–80% reduction in cryogenic system cost and complexity.

Design Constraints. Panel (f) maps the relationship between operating temperature and minimum required qubit frequency for various thermal population thresholds. The HEATS-Q operating points (blue circles) lie well above the 1% threshold curve, indicating robust operation at standard conditions, with margin for parameter drift and environmental fluctuations.




Figure 2.4: HEATS-Q operating regime analysis. **Left:** Zoomed view of thermal population as a function of temperature for qubit frequencies near the HEATS-Q sweet spot (1.972 GHz, red dashed line). The shaded regions indicate performance zones: green ($< 2\%$, excellent), light-blue (2–5%, good), and yellow (5–10%, acceptable). Operating points at 10, 12, 14, 16, and 20 mK are marked with values. **Right:** Bar chart comparing thermal population (solid bars) and normalized energy ratio (hatched bars) between traditional Al/AIO_x transmons at 20 mK and HEATS-Q at this range of temperatures. The red dashed line marks the 2% thermal population threshold.

2.4.3 Flux-Tunability and Sweet-Spot Operation

Frequency Range

The qubit frequency varies with applied flux:

$$\omega_{01}(\Phi) \approx \sqrt{8E_J(\Phi)E_C} - E_C \quad (2.31)$$

For our parameters:

- At $\Phi = 0$: $E_J = 27$ GHz $\Rightarrow \omega_{01} \approx 6.37$ GHz
- At $\Phi = \Phi_0/2$: $E_J = 3$ GHz $\Rightarrow \omega_{01} \approx 1.97$ GHz

This provides a tuning range of approximately 4.4 GHz, enabling:

1. Frequency-addressable multi-qubit architectures
2. Dynamic flux-based gates (e.g., adiabatic CZ gates)
3. Compatibility with flux-dip control schemes for high-frequency operation

Sweet-Spot Benefits

At the flux sweet spot ($\Phi = \Phi_0/2$), the first-order sensitivity to flux noise vanishes:

$$\left. \frac{\partial \omega_{01}}{\partial \Phi} \right|_{\Phi=\Phi_0/2} = 0 \quad (2.32)$$

This leads to enhanced pure dephasing time T_2^* due to suppression of low-frequency flux noise. The asymmetry $\eta = 11.1\%$ ensures:

- Non-zero $E_J(\Phi_0/2) = |E_\Delta| = 3 \text{ GHz}$ (maintaining anharmonicity)
- Reduced charge dispersion compared to symmetric SQUIDs
- Stable operating point for high-fidelity gates

2.5 Two-Qubit Cavity QED System Analysis

Physical Interpretation

The Hamiltonian consists of three parts:

1. **Cavity term** ($\hbar\omega_r\hat{a}^\dagger\hat{a}$): Free evolution of the cavity photon field
2. **Qubit terms** ($\hbar\omega_{q,i}\hat{\sigma}_z^{(i)}/2$): Free evolution of the qubits
3. **Interaction terms** ($\hbar g_i(\hat{a}^\dagger\hat{\sigma}_-^{(i)} + \hat{a}\hat{\sigma}_+^{(i)})$): Coherent energy exchange between qubits and cavity via single-photon processes

2.5.1 System Parameters for HEATS-Q

We consider the following representative parameters:

- Cavity frequency: $\omega_r/2\pi = 6.5 \text{ GHz}$
- Qubit 1 frequency: $\omega_{q1}/2\pi = 5.8 \text{ GHz}$
- Qubit 2 frequency: $\omega_{q2}/2\pi = 5.2 \text{ GHz}$
- Coupling strengths: $g_1/2\pi = g_2/2\pi = 80 \text{ MHz}$

The detunings are:

$$\Delta_1 = \omega_{q1} - \omega_r = -0.7 \text{ GHz} \quad (2.33)$$

$$\Delta_2 = \omega_{q2} - \omega_r = -1.3 \text{ GHz} \quad (2.34)$$

2.5.2 Dispersive Regime Transformation

Criterion for Dispersive Regime

The system operates in the dispersive regime when:

$$|\Delta_i| \gg g_i \quad \text{for all qubits } i \quad (2.35)$$

For our parameters:

$$\frac{|\Delta_1|}{g_1} = 8.8 \quad (\text{satisfied}) \quad (2.36)$$

$$\frac{|\Delta_2|}{g_2} = 16.2 \quad (\text{satisfied}) \quad (2.37)$$

In this regime, real photon exchange between qubits and cavity is suppressed, but virtual processes lead to effective qubit-qubit interactions mediated by the cavity.

Effective Dispersive Hamiltonian

Under a Schrieffer-Wolff transformation (second-order perturbation theory in g_i/Δ_i), the Hamiltonian becomes:

$$\hat{H}_{\text{disp}} = \hbar\tilde{\omega}_r\hat{a}^\dagger\hat{a} + \sum_{i=1}^2 \frac{\hbar\tilde{\omega}_{q,i}}{2}\hat{\sigma}_z^{(i)} + \sum_{i=1}^2 \hbar\chi_i\hat{\sigma}_z^{(i)}\hat{a}^\dagger\hat{a} + \hbar\zeta\hat{\sigma}_z^{(1)}\hat{\sigma}_z^{(2)} \quad (2.38)$$

where the dispersive shifts are:

$$\chi_i = \frac{g_i^2}{\Delta_i} \quad (2.39)$$

and the effective *transverse* qubit-qubit exchange coupling (Majer *et al.* [?]) is

$$J = \frac{g_1 g_2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right). \quad (2.40)$$

This is the coupling appearing as $J(\hat{\sigma}_+^{(1)}\hat{\sigma}_-^{(2)} + \text{h.c.})$ in the rotated frame; the residual longitudinal $\hat{\sigma}_z \otimes \hat{\sigma}_z$ cross-Kerr ζ_{zz} requires the transmon anharmonicity and is derived in Section 2.10. Note: in earlier versions of this manuscript the symbol ζ was used for the coefficient $\hbar\zeta\hat{\sigma}_z^{(1)}\hat{\sigma}_z^{(2)}$ in Eq. (2.38); this was an inadvertent mislabelling of the transverse exchange coupling, addressed in detail in Section 2.10.3.

In the dispersive regime and for weak anharmonicity, the transmon-induced dispersive shift is well approximated by

$$\frac{\chi_i}{2\pi} \simeq \frac{\left(\frac{g_i}{2\pi}\right)^2 \left(\frac{\alpha_i}{2\pi}\right)}{\left(\frac{\Delta_i}{2\pi}\right) \left(\frac{\Delta_i}{2\pi} + \frac{\alpha_i}{2\pi}\right)}, \quad (2.41)$$

where $\Delta_i = \omega_i - \omega_r$ and $\alpha_i = \omega_{12,i} - \omega_{01,i} < 0$.

The renormalized frequencies are:

$$\tilde{\omega}_r = \omega_r + \chi_1 + \chi_2 \quad (2.42)$$

$$\tilde{\omega}_{q,i} = \omega_{q,i} + \chi_i \quad (2.43)$$

Numerical Values for HEATS-Q (redesigned operating point)

Adopting the redesigned operating point $\Delta_q/2\pi = 280$ MHz ($\omega_{q,1}/2\pi = 5.640$ GHz, $\omega_{q,2}/2\pi = 5.360$ GHz; see discussion of the cross-resonance sweet spot in Section 2.11.2), the dispersive shifts become

$$\chi_1 = \frac{g_1^2}{\Delta_1} \frac{\alpha}{\Delta_1 + \alpha} = -12.9 \text{ MHz}, \quad (2.44)$$

$$\chi_2 = \frac{g_2^2}{\Delta_2} \frac{\alpha}{\Delta_2 + \alpha} = 24.3 \text{ MHz}, \quad (2.45)$$

and the transverse exchange coupling is

$$J = \frac{g_1 g_2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right) = -6.5 \text{ MHz}. \quad (2.46)$$

The corresponding longitudinal cross-Kerr from Eq. (2.109) is $\zeta_{zz}/2\pi \simeq -1.7$ MHz: large enough to require active cancellation via the echo-CR sequence (Section 2.11.3), but

suppressed relative to J by a factor ~ 4 . The chosen working point sits on the high- Δ_q side of the straddling resonance $\Delta_q = |\alpha| = 200$ MHz; the actual ζ_{zz} value should be calibrated experimentally and the operating point fine-tuned accordingly.

Erratum to previous numerical values. The pre-revision version of this section reported $\chi_1 = -9.14$ MHz, $\chi_2 = -4.92$ MHz and a single coupling $\zeta = -7.03$ MHz, the latter quoted as the $\sigma_z \otimes \sigma_z$ coupling. As explained above (and in Section 2.10.3), the formula $(g_1 g_2 / 2)(\Delta_1^{-1} + \Delta_2^{-1})$ in fact gives the transverse exchange J , not the longitudinal cross-Kerr ζ_{zz} , and the operating point has been redesigned to $\Delta_q = 280$ MHz to enable a fast, high-fidelity cross-resonance gate.

The frequency differences are:

$$\omega_r(|00\rangle) - \omega_r(|01\rangle) = 2\chi_2 = 9.84 \text{ MHz} \quad (2.47)$$

$$\omega_r(|00\rangle) - \omega_r(|10\rangle) = 2\chi_1 = 18.28 \text{ MHz} \quad (2.48)$$

$$\omega_r(|00\rangle) - \omega_r(|11\rangle) = 2(\chi_1 + \chi_2) = 28.12 \text{ MHz} \quad (2.49)$$

Readout Protocol

1. Drive cavity with weak coherent tone at ω_r
2. Measure transmitted/reflected signal using IQ demodulation
3. Phase and amplitude depend on qubit state
4. Discriminate states based on measured IQ point

Resolution Criterion

For good state discrimination, the dispersive shift must exceed the cavity linewidth:

$$2|\chi_i| > \kappa \quad (2.50)$$

Assuming $\kappa/2\pi = 2$ MHz:

$$\frac{2|\chi_1|}{\kappa} = \frac{18.28}{2} = 9.14 \quad (\text{excellent}) \quad (2.51)$$

$$\frac{2|\chi_2|}{\kappa} = \frac{9.84}{2} = 4.92 \quad (\text{good}) \quad (2.52)$$

Both conditions are satisfied, ensuring high-fidelity readout.

2.5.3 Single-Qubit Gates

Single-qubit rotations are implemented with resonant microwave drives at the qubit frequency. A general rotation is:

$$R_i(\theta, \phi) = \exp\left(-i\frac{\theta}{2} \left[\cos(\phi)\hat{\sigma}_x^{(i)} + \sin(\phi)\hat{\sigma}_y^{(i)}\right]\right) \quad (2.53)$$

Common gates:

$$X_i = R_i(\pi, 0) \quad (\text{bit flip}) \quad (2.54)$$

$$Y_i = R_i(\pi, \pi/2) \quad (2.55)$$

$$H_i = R_i(\pi/2, \pi/2) \cdot R_i(\pi, 0) \quad (\text{Hadamard}) \quad (2.56)$$

Typical gate times are $\tau_{1Q} \sim 20$ ns for high-fidelity operations with DRAG pulse shaping.

2.5.4 Two-Qubit Gates: CNOT Gate Implementation

CNOT Truth Table

The CNOT gate with qubit 1 as control and qubit 2 as target produces:

Input	Output
$ 00\rangle$	$ 00\rangle$
$ 01\rangle$	$ 01\rangle$
$ 10\rangle$	$ 11\rangle$
$ 11\rangle$	$ 10\rangle$

(2.57)

This implements the logical operation: if control qubit is $|1\rangle$, flip the target qubit.

CNOT Implementation Strategy: CZ Gate + Single-Qubit Rotations

The CNOT gate can be decomposed as:

$$\text{CNOT}_{12} = R_y^{(2)}\left(-\frac{\pi}{2}\right) \cdot \text{CZ} \cdot R_y^{(2)}\left(\frac{\pi}{2}\right) \quad (2.58)$$

where the controlled-Z (CZ) gate is:

$$\text{CZ} = \exp\left(-i\frac{\pi}{4}\hat{\sigma}_z^{(1)}\hat{\sigma}_z^{(2)}\right) \quad (2.59)$$

2.5.5 Method 1: $\sqrt{i\text{SWAP}}$ Gate (Resonant Interaction)

Physical Mechanism

The $\sqrt{i\text{SWAP}}$ approach uses *resonant* coupling between the two qubits. The key steps are:

1. Apply flux pulses to tune qubit frequencies into resonance: $\omega_{q1} \approx \omega_{q2}$
2. Strong exchange interaction at rate $g_{\text{eff}} \approx g \sim 80$ MHz
3. Partial swap operation for time $t = \pi/(4g_{\text{eff}})$
4. Return qubits to original frequencies
5. Apply single-qubit corrections

Gate Time Calculation

For a single $\sqrt{i\text{SWAP}}$ operation:

$$t_{\sqrt{i\text{SWAP}}} = \frac{\pi}{4g_{\text{eff}}} = \frac{\pi}{4 \times 80 \text{ MHz}} \approx 9.8 \text{ ns} \quad (2.60)$$

To implement a CNOT gate, we need:

$$t_{\text{CNOT}} = 2t_{\sqrt{i\text{SWAP}}} + t_{\text{rotations}} \approx 20 + 40 + 20 = \mathbf{80 \text{ ns}} \quad (2.61)$$

where the additional time accounts for single-qubit rotations and flux pulse overhead.

Advantages and Disadvantages

Advantages:

- Fast gate time (80 ns)
- High coupling strength
- Mature technology (used by Google, Rigetti)
- Experimentally demonstrated with high fidelity

Disadvantages:

- Requires flux control lines for frequency tuning
- More complex pulse sequences
- Potential leakage to $|2\rangle$ state when near resonance
- Sensitive to frequency fluctuations
- Less robust at elevated temperatures

2.5.6 Method 2: ZZ-Coupling Gate (Dispersive Interaction)

Physical Mechanism

The ZZ-coupling approach uses the *always-on* dispersive interaction between qubits mediated by the cavity. No frequency tuning is required.

The effective Hamiltonian in the dispersive limit is:

$$H_{\text{eff}} = \sum_i \tilde{\omega}_i \frac{\sigma_z^{(i)}}{2} + J_{zz} \sigma_z^{(1)} \otimes \sigma_z^{(2)} \quad (2.62)$$

where the cavity-mediated ZZ coupling is calculated using the dispersive Hamiltonian approach [103, 111]:

$$J_{zz} = \frac{g_1 g_2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right) \quad (2.63)$$

With our parameters:

$$J_{zz} = \frac{80 \times 80}{2} \times \left(\frac{1}{-700} + \frac{1}{-1300} \right) \text{ MHz} \quad (2.64)$$

$$= 6.4 \times (-2.198 \times 10^{-3}) \text{ MHz} \quad (2.65)$$

$$= -7.03 \text{ MHz} \quad (2.66)$$

A controlled-phase (CZ) gate requires a conditional phase of $\pi/2$:

$$t_{\text{CZ}} = \frac{\pi}{2|J_{zz}|} = \frac{\pi}{2 \times 7.03 \text{ MHz}} \approx \mathbf{224 \text{ ns}} \quad (2.67)$$

The CZ gate is equivalent to CNOT up to single-qubit rotations.

2.5.7 Physical Origin of Slower Gate

The key difference is the coupling strength. The ZZ coupling arises from second-order perturbation theory:

$$J_{zz} \sim \frac{g^2}{\Delta} \quad (2.68)$$

Since $g/\Delta \ll 1$ (dispersive regime), we have:

$$\frac{J_{zz}}{g} \sim \frac{g}{\Delta} \approx \frac{80}{700} \approx 0.11 \quad (2.69)$$

The effective coupling is weaker by approximately one order of magnitude, leading to proportionally longer gate times:

$$\frac{t_{ZZ}}{t_{\sqrt{i}\text{SWAP}}} \approx \frac{g}{J_{zz}} \approx \frac{80}{14} \approx 5.7 \quad (2.70)$$

Numerically: $224/80 \approx 2.8$ (including overhead from different protocols).

Advantages and Disadvantages

Advantages:

- **No flux control required** – qubits remain at fixed frequencies
- Simpler hardware (fewer control lines)
- No leakage to higher levels (always far detuned)
- More robust to frequency fluctuations
- Better suited for elevated-temperature operation
- Easier to scale to many qubits

Disadvantages:

- Slower gate time (224 ns vs 80 ns)
- Weaker coupling strength

Performance Comparison

Fidelity Estimates

Assuming decoherence-limited fidelity with $T_2 = 50 \mu\text{s}$ (achievable with echo sequences at 4–6 K):

$$F_{\text{CNOT}}^{\sqrt{i}\text{SWAP}} \approx 1 - \frac{t_{\text{gate}}}{T_2} = 1 - \frac{80}{50000} = \text{99.84\%} \quad (2.71)$$

$$F_{\text{CNOT}}^{\text{ZZ}} \approx 1 - \frac{t_{\text{gate}}}{T_2} = 1 - \frac{224}{50000} = \text{99.55\%} \quad (2.72)$$

The difference is only $\Delta F = 0.06\%$ – both well above the fault-tolerance threshold of 99%.

2.5.8 Recommendation for HEATS-Q

For the HEATS-Q project operating at standard mk temperatures, we recommend the **ZZ-coupling approach (224 ns)** for the following reasons:

Table 2.3: Comparison of Two-Qubit Gate Implementations

Property	$\sqrt{i}\text{SWAP}$	ZZ Coupling
Coupling strength	$g \sim 80$ MHz	$J_{zz} \sim 14$ MHz
Gate time	80 ns	224 ns
Estimated fidelity	99.84%	99.55%
Flux control	Required	Not required
Frequency tuning	Yes (dynamic)	No (fixed)
Hardware complexity	High	Low
Leakage risk	Higher	Lower
Calibration	Complex	Simple
Noise sensitivity	Higher	Lower
Best for	Maximum speed	Robustness

Performance

Despite the 144 ns longer gate time:

- Gate time (224 ns) is a bit over the 200 ns “fast gate” threshold
- Fidelity (99.55%) exceeds the 99% surface code threshold
- The 0.29% fidelity difference is negligible in practice

Scalability

For multi-qubit processors:

- Fewer control lines required (N microwave lines vs N microwave + N flux lines)
- Lower crosstalk between qubits
- Simpler cryogenic wiring at elevated temperatures

Final considerations

Both reported gate times are potentially valid choices:

- **80 ns:** Corresponds to $\sqrt{i}\text{SWAP}$ implementation using resonant coupling
- **224 ns:** Corresponds to ZZ-coupling implementation using dispersive interaction

The final choice between methods represents a trade-off between speed and simplicity. For HEATS-Q, the ZZ-coupling approach could be preferred because:

1. Compatible with flux-bias-free asymmetric SQUID design
2. More robust at elevated temperatures (4–6 K)
3. Simpler hardware and calibration
4. Achieves excellent performance (99.55% fidelity)
5. Better suited for scalable architectures

The 30 ns increase in gate time is a small price to pay for significantly reduced system complexity and improved robustness – essential considerations for practical quantum computing at elevated temperatures.

This analysis demonstrates that both gate implementations are valid and have been successfully demonstrated in state-of-the-art superconducting quantum processors. The $\sqrt{i}$ SWAP approach is used by Google [12] and Rigetti, while the ZZ-coupling approach has been extensively studied by IBM [103, 111] and others. The choice between them depends on system priorities: speed vs. simplicity and robustness.

Operation	Time	Fidelity	Method
Initialization	$5T_1$	$> 99.9\%$	Thermal + active reset
Single-qubit gates	20 ns	$> 99.9\%$	Resonant MW drive
Readout	200 ns	$> 99\%$	Dispersive (homodyne)
CNOT (ZZ coupling)	225 ns	$> 99\%$	practical
CNOT (resonant swap)	80 ns	$> 99\%$	Flux-tunable coupling

Table 2.4: Summary of quantum operation metrics for two-qubit cavity QED system.

2.6 Two-Qubit System Numerical Analysis

In this section, we present the numerical analysis summary of two transmon qubits coupled through a microwave cavity resonator in the dispersive regime, optimized for operation at 10 mK to 20 mK temperatures.

2.6.1 System Architecture

Figure 2.5 shows the physical architecture of the two-qubit system. Two transmon qubits (Q1 and Q2) are coupled to a common cavity resonator with coupling strengths g_1 and g_2 . The qubits are detuned from the cavity frequency, placing the system in the dispersive regime where virtual photon exchange mediates an effective qubit-qubit interaction with strength $J_{zz} = 7.03$ MHz.

2.6.2 System Parameters

- Cavity frequency: $\omega_r/2\pi = 6.5$ GHz
- Qubit 1 frequency: $\omega_{q1}/2\pi = 5.8$ GHz (detuning: $\Delta_1 = -0.7$ GHz)
- Qubit 2 frequency: $\omega_{q2}/2\pi = 5.2$ GHz (detuning: $\Delta_2 = -1.3$ GHz)
- Coupling strengths: $g_1 = g_2 = 80$ MHz
- Dispersive shifts: $\chi_1 = -2.03$ MHz, $\chi_2 = -0.66$ MHz
- ZZ coupling: $J_{zz} = 7.03$ MHz
- CNOT gate time: $t_{CNOT} = 224$ ns
- Gate fidelity: $F = 99.55\%$

2.6.3 Energy Level Structure

The energy level structure is shown in Figure 2.7. The two qubits have different transition frequencies (5.8 GHz and 5.2 GHz), allowing individual addressability. Both qubits are detuned below the cavity frequency (6.5 GHz) by $\Delta_1 = -0.7$ GHz and $\Delta_2 = -1.3$ GHz,




Figure 2.5: **Two-qubit cavity QED system architecture.** Two transmon qubits (Q1 at 5.8 GHz, Q2 at 5.2 GHz) couple to a shared cavity resonator at 6.5 GHz with coupling strengths $g_1 = g_2 = 80 \text{ MHz}$. The qubits are red-detuned from the cavity ($\Delta_1 = -0.7 \text{ GHz}$, $\Delta_2 = -1.3 \text{ GHz}$), enabling dispersive operation. Virtual photon exchange through the cavity creates an effective ZZ coupling of 7.03 MHz between the qubits, which enables two-qubit gates without direct interaction.

respectively. These large detunings ($|\Delta|/g \gg 1$) ensure the system operates in the dispersive regime where energy exchange is suppressed, and the dynamics are governed by effective frequency shifts.

2.6.4 Dispersive Frequency Shifts and Readout

Figure 2.8 demonstrates the dispersive frequency shifts that enable qubit state readout. Panel (a) shows how the cavity frequency shifts by $\pm\chi_i$ depending on the state of each individual qubit. Panel (b) shows that all four two-qubit states produce distinct cavity frequencies, separated by at least 1.31 MHz, which exceeds the typical cavity linewidth of 1 MHz and enables high-fidelity state discrimination.




Figure 2.6: **Two-qubit cavity QED system architecture and parameters.** This baseline architecture demonstrates state-of-the-art performance at standard millikelvin temperatures before extension to elevated-temperature operation. **Left panel:** Physical layout showing two transmon qubits (Q1 at 5.8 GHz, Q2 at 5.2 GHz) coupled to a shared cavity resonator (6.5 GHz) with coupling strengths $g_1 = g_2 = 80$ MHz. Detunings $\Delta_1 = -0.7$ GHz and $\Delta_2 = -1.3$ GHz place the system in the dispersive regime ($|\Delta_i|/g_i > 5$). Virtual photon exchange mediates effective ZZ coupling $\zeta_{12}/2\pi = 7.03$ MHz (purple dashed line). **Right panel:** Complete system parameters including cavity quality factor $Q = 6500$, anharmonicities $\alpha_1/2\pi = -0.70$ GHz and $\alpha_2/2\pi = -1.30$ GHz, and dispersive shifts $\chi_1/2\pi = -2.03$ MHz and $\chi_2/2\pi = -0.66$ MHz, enabling high-fidelity quantum non-demolition readout with $\chi_i/\kappa \gg 1$.




Figure 2.7: **Energy level structure and detunings.** Energy levels of the two qubits and cavity mode. Qubit 1 (red, 5.8 GHz) and Qubit 2 (green, 5.2 GHz) are both red-detuned from the cavity (blue, 6.5 GHz). The detunings $\Delta_1 = -0.7$ GHz and $\Delta_2 = -1.3$ GHz satisfy the dispersive regime condition $|\Delta_i|/g_i > 5$, with ratios of 8.8 and 16.2 respectively. This large detuning suppresses direct photon exchange while enabling dispersive interactions.




Figure 2.8: **Dispersive frequency shifts and two-qubit state readout.** (a) Single-qubit dispersive shifts. The cavity frequency shifts by $\chi_1 = -2.03$ MHz for Qubit 1 and $\chi_2 = -0.66$ MHz for Qubit 2, with opposite signs for $|0\rangle$ and $|1\rangle$ states. (b) Two-qubit state readout. Each of the four computational basis states ($|00\rangle$, $|01\rangle$, $|10\rangle$, $|11\rangle$) produces a unique cavity frequency according to $\omega_c = \omega_r + \chi_1 s_{z1} + \chi_2 s_{z2}$, where $s_z = \pm 1$. The minimum frequency separation of 1.31 MHz exceeds the cavity linewidth ($\kappa \sim 1$ MHz), enabling single-shot readout of all two-qubit states with high fidelity.

2.6.5 ZZ Coupling and Gate Operations

Figure 2.9 characterizes the effective qubit-qubit interaction and gate performance. Panel (a) shows the ZZ coupling energy spectrum, where states $|00\rangle$ and $|11\rangle$ have the same energy shift (-7.03 MHz), and states $|01\rangle$ and $|10\rangle$ have the opposite shift ($+7.03$ MHz). This pattern of conditional energy shifts enables entangling operations. Panel (b) compares the timing of various quantum operations, showing that the CNOT gate time of 224 ns is still fast enough to maintain high fidelity while remaining simple to implement without flux control.

The system operates without flux control, using only fixed-frequency qubits and microwave pulses, making it well-suited for elevated-temperature operation at 4–6 K where thermal stability is important.




Figure 2.9: **ZZ coupling spectrum and gate timing.** (a) ZZ coupling energy spectrum. The cavity-mediated interaction produces an effective Hamiltonian $H_{ZZ} = J_{zz}\sigma_z^{(1)} \otimes \sigma_z^{(2)}$ with coupling strength $J_{zz} = 7.03$ MHz, calculated using the formula $J_{zz} = (g_1 g_2 / 2) (\frac{1}{\Delta_1} + \frac{1}{\Delta_2})$ from Gambetta et al. (2011). States where the qubits are aligned ($|00\rangle$, $|11\rangle$) have energy shift $-J_{zz}$, while anti-aligned states ($|01\rangle$, $|10\rangle$) have shift $+J_{zz}$. (b) Gate timing comparison. Single-qubit gates (20 ns) and the CNOT gate (224 ns) are a bit over the 200 ns threshold for fast gates. The CNOT time $t = \pi / (2|J_{zz}|)$ corresponds to the conditional phase evolution required for an entangling gate. Readout (300 ns) includes cavity ring-up and signal integration.

2.6.6 Performance Summary

The two-qubit system demonstrates excellent performance for fault-tolerant quantum computing:

- **Dispersive regime:** $|\Delta_1|/g_1 = 8.8$ and $|\Delta_2|/g_2 = 16.2$ (requirement: > 5)
- **ZZ coupling:** $J_{zz} = 7.03$ MHz (formula verified against Gambetta et al., 2011)
- **CNOT gate time:** 224 ns (a bit over 200 ns fast gate threshold)
- **Gate fidelity:** $F = \exp(-t_{\text{gate}}/T_2) = 99.55\%$ with $T_2 = 50$ μs
- **Error rate:** $\epsilon = 0.45\%$ (below 1% surface code threshold)

- **Readout:** All four two-qubit states distinguishable (minimum separation 1.31 MHz)

The system operates without flux control, using only fixed-frequency qubits and microwave pulses, making it well-suited for elevated-temperature operation at 4–6 K where thermal stability is important.

The CNOT gate time is then:

$$t_{\text{CNOT}} = \frac{\pi}{2|J_{zz}|} = \frac{\pi}{2 \times 7.03 \text{ MHz}} \approx 224 \text{ ns} \quad (2.73)$$

2.6.7 Two-Qubit Operation Take-away

The four figures presented demonstrate that this two-qubit cavity QED system achieves:

1. Clean dispersive operation with large detuning ratios
2. State-dependent cavity shifts enabling high-fidelity readout
3. Distinguishability of all four two-qubit computational basis states
4. Strong effective ZZ coupling (7.03) MHz for fast entangling gates
5. CNOT gate performance (224 ns, 99.55% fidelity) exceeding fault-tolerance requirements

These results validate the design for two-qubit quantum operations at elevated temperatures, with performance parameters consistent with state-of-the-art superconducting quantum processors operating at millikelvin temperatures.

Proven Baseline

Baseline Performance Summary at 10–20 mK:

- Single-qubit gates: $F > 99.9\%$
- Two-qubit CNOT: $F = 99.55\%$
- Thermal effects: Negligible ($< 10^{-5}\%$)
- Status: **Competitive with state-of-the-art**

2.7 Baseline Validation

Our baseline analysis demonstrates that asymmetric SQUID qubits at 5–6 GHz operating at 10–20 mK achieve:

1. **High gate fidelities:** Single-qubit $> 99.9\%$, two-qubit $> 99.5\%$
2. **Fast gates:** 20 ns single-qubit, 224 ns two-qubit
3. **Robust dispersive operation:** $|\Delta|/g > 8$
4. **Negligible thermal noise:** $n_{\text{th}} < 10^{-5}$

Conclusion: The baseline design is sound and achieves performance comparable to leading quantum processors.

Part C

Elevated-Temperature Operation at 300 GHz

2.8 High-frequency, high-temperature environment system

Standard quantum processors operating at 10–20 mK face significant challenges:

- **Cooling cost:** Dilution refrigerators cost \$500k–\$2M
- **Cooling time:** Days to reach operating temperature
- **Power consumption:** kW-scale at room temperature
- **Scalability:** Difficult to scale beyond hundreds of qubits
- **Integration:** Cannot co-locate with classical electronics

Operating at 4–10 K would enable:

- **Pulse-tube coolers:** \$50k–\$100k (10× cheaper)
- **Fast cooldown:** Hours instead of days
- **Scalability:** Modular, distributed systems
- **Integration:** Co-locate with cryo-CMOS control

The Challenge: Thermal energy at 4 K is much higher. We need a radically different approach.

Thermal Energy at Elevated Temperatures

At temperature T , the thermal energy scale is:

$$\frac{k_B T}{h} = 20.84 \text{ GHz/K} \quad (2.74)$$

Therefore:

$$\text{At } T = 4 \text{ K : } \frac{k_B T}{h} = 83.3 \text{ GHz} \quad (2.75)$$

$$\text{At } T = 10 \text{ K : } \frac{k_B T}{h} = 208.4 \text{ GHz} \quad (2.76)$$

Quantum Regime Condition

The quantum ratio at 300 GHz and 4 K is:

$$R_Q = \frac{\hbar\omega_{q1}}{k_B T} = \frac{300 \text{ GHz}}{83.3 \text{ GHz}} = 3.60 \quad (2.77)$$

This value satisfies $R_Q > 3$, placing the system in a good quantum regime. Figure 2.12(a) shows the quantum ratio landscape across temperature and frequency space.

Thermal Population

The thermal occupation of the first excited state is:

$$n_{\text{th}} = \frac{1}{\exp(R_Q) - 1} = \frac{1}{\exp(3.60) - 1} = 0.0281 = 2.81\% \quad (2.78)$$

This thermal population contributes to gate errors but remains manageable. At lower temperatures, n_{th} decreases exponentially:

$$\text{At } T = 3 \text{ K : } n_{\text{th}} = 0.56\% \quad (2.79)$$

$$\text{At } T = 2 \text{ K : } n_{\text{th}} = 0.07\% \quad (2.80)$$

Figure 2.12(b) shows thermal population versus temperature at 300 GHz.

Design Requirements

The coupling strength g must balance two competing requirements:

1. **Strong enough** for fast two-qubit gates: $J_{zz} \propto g^2/\Delta$
2. **Weak enough** for valid dispersive regime: $|\Delta|/g > 5$

Comparison with conventional systems:

- **5.8 GHz at 20 mK:** $\hbar\omega/(k_B T) = 14.0$, $n_{\text{th}} = 0.08\%$
- **300 GHz at 4 K:** $\hbar\omega/(k_B T) = 3.60$, $n_{\text{th}} = 2.81\%$ ✓

Temperature	$\hbar\omega/(k_B T)$	n_{th}	Assessment
20 mK	720	1.4×10^{-312}	Excellent
2 K	7.20	0.075%	Excellent
4 K	3.60	2.81%	Good
6 K	2.40	9.98%	Acceptable
8 K	1.80	19.8%	Marginal
10 K	1.44	31.1%	Challenging

Table 2.5: Thermal populations at 300 GHz across temperatures

Novel Innovation

Key Innovation:

By operating at 300 GHz, we maintain quantum regime operation at elevated temperatures:

- At 4 K: $n_{\text{th}} = 2.8\%$ (manageable)
- At 6 K: $n_{\text{th}} = 10\%$ (acceptable with error correction)
- Compare to 5.8 GHz at 4 K: $n_{\text{th}} = 1388\%$ (impossible!)

Configuration	Frequency	Temperature	n_{th}	Status
Baseline	5.8 GHz	20 mK	1.4×10^{-5}	✓ Proven
Baseline at 4 K	5.8 GHz	4 K	1388%	× Impossible
Innovation	300 GHz	4 K	2.8%	Novel
Innovation	300 GHz	6 K	10%	Challenging

Table 2.6: Baseline vs innovative approach

This six-panel analysis of Fig.2.11 provides a complete characterization of thermal behavior for high-frequency qubits at elevated temperatures. **Panel (a):** Thermal population versus temperature for different qubit frequencies (100–400 GHz). Logarithmic scale reveals that 300 GHz operation (red line) maintains $n_{\text{th}} < 5\%$ up to 6 K (green region) and crosses the 10% acceptable threshold near 8 K (yellow region). The 300 GHz-Q operating range (4–10 K, shaded area) is highlighted. Lower frequencies (<200 GHz) exhibit prohibitively high thermal populations even at 4 K, while higher frequencies (>350 GHz) offer marginal improvement at the cost of increased fabrication and control complexity. **Panel (b):** Thermal population versus qubit frequency at different temperatures (4–20 K). At 4 K (bottom blue curve), frequencies above 200 GHz achieve quantum regime operation ($n_{\text{th}} < 10\%$). The red star marks the 300 GHz "typical transmons" sweet spot at 4 K ($n_{\text{th}} = 2.81\%$). Classical-to-quantum crossover occurs around 100 GHz at 4 K and 250 GHz at 10 K, with the 300 GHz-Q range (vertical purple band) providing a robust operating window. **Panel (c):** Two-dimensional thermal population landscape (logarithmic color scale). The 300 GHz-Q sweet spot design (red star) sits comfortably in a low-thermal-excitation zone (blue region). The landscape reveals an exponential suppression boundary (black dashed contour at 5% threshold) that defines the feasible operating region. White dashed box indicates the 300 GHz-Q parameter space (250–350 GHz, 4–10 K). **Panel (d):** Energy ratio $\hbar\omega/(k_B T)$ versus temperature for different frequencies, with performance thresholds indicated: green (ratio > 10, excellent), yellow (ratio > 5, good), orange (ratio > 3, acceptable), and red (ratio > 1, poor). The 300 GHz design maintains ratio > 3 across the entire 4–8 K window, ensuring quantum behavior dominates thermal fluctuations. Traditional transmons achieve ratios $\sim 10^3$ but require dilution refrigeration. **Panel (e):** Performance comparison table quantifying thermal populations n_{th} , energy ratios, and status assessment for traditional Al-based transmons (5 GHz at 20 mK) versus 300 GHz-Q operation at 4, 6, 8, and 10 K. Traditional configuration achieves essentially zero thermal noise (0.020%, excellent) with energy ratio 11998. The 300 GHz-Q at 4 K maintains good performance (2.81%, ratio 3.60), transitioning to marginal at 6 K (9.98%, ratio 2.40) and poor beyond 8 K (19.81%, ratio 1.80).




Figure 2.10: **Operating regime analysis for 300 GHz qubits at elevated temperatures.** **Left panel:** Thermal population percentage versus temperature for different qubit frequencies near the 300 GHz sweet spot. The black dashed line at $\omega_{01} = 300$ GHz (target design point) shows thermal population increasing from 2.81% at 4 K to 31.06% at 10 K. Shaded regions indicate performance zones: green (excellent, $n_{th} < 2\%$), light blue (good, 2–5%), and yellow (acceptable, 5–10%). Operating points at 4, 5, 6, 8, and 10 K are marked with quantitative thermal population values, demonstrating that the system maintains acceptable thermal noise up to ~ 6 K. Frequencies below 250 GHz exhibit unacceptably high thermal populations ($> 10\%$) even at 4 K, while frequencies above 350 GHz offer diminishing returns. The 300 GHz choice represents an optimal balance between thermal suppression and experimental accessibility. **Right panel:** Performance comparison of traditional Al/AIO_x transmons at 5.8 GHz/20 mK versus 300 GHz-Q operation at elevated temperatures. Solid blue bars show thermal population percentage (left axis), while hatched orange bars display normalized energy ratio $\hbar\omega/(k_B T)/20$ (right axis, scaled for visibility). The traditional configuration achieves essentially zero thermal population ($< 0.0001\%$, imperceptible on scale) and extreme energy ratio 0.7 (representing $\hbar\omega/(k_B T) \approx 14$). In stark contrast, 300 GHz-Q at 4 K maintains manageable 2.81% thermal population with energy ratio 0.2 (representing $\hbar\omega/(k_B T) = 3.60$, satisfying quantum regime criterion > 3). At 6 K and 8 K, thermal populations rise to 9.98% and 19.81% respectively, approaching but remaining within acceptable bounds for quantum error correction ($< 20\%$). The red dashed horizontal line at 2% marks the excellent performance threshold. This analysis demonstrates the feasibility of elevated-temperature quantum computing through high-frequency operation, trading five orders of magnitude in thermal suppression for 200–500 $\times$ higher operating temperature, with corresponding 70–80% reduction in cryogenic infrastructure cost and complexity.




Figure 2.11: Comprehensive thermal analysis of 300 GHz qubit operation across parameter space.

This confirms the 4–6 K operating window for high-fidelity quantum operations. **Panel (f)**: Minimum required qubit frequency versus temperature to maintain specified thermal population thresholds (1%, 2%, 5%, 10%). The 300 GHz-Q design points (blue circles) lie well above the 2% curve and comfortably above the 5% threshold, providing margin for parameter drift and environmental fluctuations. The traditional 20 mK transmon at 5 GHz (green square at left edge) achieves $< 1\%$ thermal population despite low frequency due to extreme temperature suppression. To extend 300 GHz-Q operation to 15 K while maintaining $< 2\%$ thermal population would require increasing frequency to ~ 450 GHz (red dashed line extrapolation). Overall, this comprehensive analysis establishes that 300 GHz operation at 4–6 K represents an optimal design point, balancing quantum regime requirements ($\hbar\omega \gg k_B T$) against experimental feasibility, enabling a transformative reduction in cryogenic complexity while maintaining performance compatible with fault-tolerant quantum error correction.

CNOT Gate Time Analysis

For a CNOT gate implemented via ZZ coupling, the gate time is:

$$t_{\text{CNOT}} = \frac{\pi}{2|J_{zz}|} \quad (2.81)$$

where the ZZ coupling strength is approximately:

$$J_{zz} \approx \frac{g^2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right) \quad (2.82)$$

For a target gate time of $t_{\text{CNOT}} = 150$ ns, we require:

$$J_{zz} = \frac{\pi}{2 \times 150 \text{ ns}} = 10.5 \text{ MHz} \quad (2.83)$$

Optimal Coupling Value

With detunings $\Delta_1 = -20$ GHz and $\Delta_2 = -40$ GHz, we estimate:

$$g \approx \sqrt{2\Delta_{\text{avg}} \cdot J_{zz}} = \sqrt{2 \times 30 \text{ GHz} \times 10.5 \text{ MHz}} = 793 \text{ MHz} \quad (2.84)$$

However, we adopt a more conservative value of $g = 500$ MHz to ensure robust operation in the dispersive regime, yielding $J_{zz} = 9.4$ MHz and $t_{\text{CNOT}} = 168$ ns.

Detuning Selection

The detunings are:

$$\Delta_1 = \omega_{q1} - \omega_r = 300 - 320 = -20 \text{ GHz} \quad (2.85)$$

$$\Delta_2 = \omega_{q2} - \omega_r = 280 - 320 = -40 \text{ GHz} \quad (2.86)$$

The dispersive ratios are:

$$\frac{|\Delta_1|}{g} = \frac{20 \text{ GHz}}{500 \text{ MHz}} = 40 \gg 5 \quad \checkmark \quad (2.87)$$

$$\frac{|\Delta_2|}{g} = \frac{40 \text{ GHz}}{500 \text{ MHz}} = 80 \gg 5 \quad \checkmark \quad (2.88)$$

Both ratios greatly exceed the requirement $|\Delta|/g > 5$, confirming valid dispersive operation. Large detunings also suppress Purcell decay (Section 2.11.6).

Energy Scales

For a transmon qubit with $E_J/E_C = 50$, the qubit frequency is:

$$\omega_q = \frac{\sqrt{8E_J E_C}}{\hbar} \quad (2.89)$$

With $E_J = 50E_C$:

$$\omega_q = \frac{\sqrt{8 \cdot 50E_C^2}}{\hbar} = \frac{20E_C}{\hbar} \quad (2.90)$$

Therefore:

$$E_C = \frac{\hbar\omega_q}{20} = \frac{(1.055 \times 10^{-34})(2\pi \times 300 \times 10^9)}{20} = \boxed{9.94 \times 10^{-24} \text{ J}} \quad (2.91)$$

$$E_C = \boxed{0.062 \text{ meV} = 15.0 \text{ GHz}} \quad (2.92)$$

$$E_J = 50E_C = \boxed{3.10 \text{ meV} = 750 \text{ GHz}} \quad (2.93)$$

The anharmonicity is:

$$\alpha = -\frac{E_C}{\hbar} = \boxed{-15.0 \text{ GHz}} \quad (2.94)$$

300 GHz / 4K System Parameters




Figure 2.12: **System Parameters and Thermal Analysis.** (a) Quantum ratio $\hbar\omega/k_B T$ versus temperature for various frequencies. The 300 GHz operating point at 4 K (orange star) achieves $R_Q = 3.6$, placing it in the good quantum regime (green region). For comparison, the baseline 5.8 GHz system (green circle) operates at 20 mK with $R_Q = 13.9$. Reference lines show excellent ($R_Q > 5$, green), good ($R_Q = 3-5$, yellow), and poor ($R_Q < 3$, red) quantum regimes. (b) Thermal population n_{th} at 300 GHz versus temperature. At the target operating point of 4 K (highlighted), thermal population is 2.81%, manageable for quantum operations. Green and orange dashed lines indicate 5% and 10% thresholds. (c) Josephson junction diameter versus current density J_c for required critical current $I_c = 30.2 \mu A$. Higher current densities (1000-2000 A/cm^2 , green region) yield feasible junction sizes (620-880 nm). Lower densities produce impractically large junctions. (d) Dispersive parameters: dispersive shifts $|\chi_i|$ and ZZ coupling J_{zz} (left axis, blue/cyan bars), and dispersive ratios $|\Delta_i|/g$ (right axis, green bars). All values support quantum non-demolition readout and two-qubit gates. (e) Transmon energy level diagram showing anharmonicity $\alpha = -1$ GHz. Energy spacings deviate from harmonic oscillator (dashed), enabling selective addressing of $|0\rangle \leftrightarrow |1\rangle$ transition. Thermal populations at 4 K are indicated (left side).

Parameter	Baseline	Innovation (300 GHz)
Qubit 1 frequency	5.8 GHz	300 GHz
Qubit 2 frequency	5.2 GHz	280 GHz
Cavity frequency	6.5 GHz	320 GHz
Coupling strength	80 MHz	500 MHz
Detuning	0.7–1.3 GHz	10–20 GHz
Operating temperature	10–20 mK	4–10 K

Table 2.7: Baseline vs 300 GHz system parameters

Critical Current

The critical current is:

$$I_c = \frac{2\pi E_J}{\Phi_0} = \frac{2\pi(4.97 \times 10^{-22})}{2.07 \times 10^{-15}} = \boxed{1.51 \mu\text{A}} \quad (2.95)$$

Junction Geometry: Area and Diameter Calculation

The junction area depends on the current density J_c :

$$A = \frac{I_c}{J_c} \quad (2.96)$$

For circular junctions, the diameter is:

$$d = \sqrt{\frac{4A}{\pi}} = \sqrt{\frac{4I_c}{\pi J_c}} \quad (2.97)$$

Table 2.8 shows junction dimensions for various current densities.

Table 2.8: Junction dimensions versus current density for $I_c = 30.2 \mu\text{A}$

J_c (A/cm ²)	Area (μm ²)	Diameter (nm)	Feasibility
100	1.510	1387	Too large ($> 1 \mu\text{m}$)
500	0.302	620	Feasible but large
1000	0.151	439	Challenging
2000	0.075	310	Challenging
5000	0.030	196	Very challenging

Critical Finding: Single junctions at 300 GHz with $E_J/E_C = 50$ require:

- $J_c \geq 3000 \text{ A/cm}^2$ for $d < 300 \text{ nm}$
- Ultra-thin AlN_x barriers ($< 0.8 \text{ nm}$) [79]
- At the edge of current fabrication capability

Optimal Design Point

We select $J_c = 1000 - 1500 \text{ A/cm}^2$, yielding junction diameters of 620–760 nm. This range:

- is achievable with modern e-beam lithography (10-20 nm resolution)
- Provides margin for fabrication tolerances
- Enables flux tunability through SQUID loop
- Allows junction arrays if needed for current distribution

Why Junction Arrays are Necessary

The fundamental problem: maintaining $E_J/E_C = 50$ at 300 GHz requires $I_c = 1.51 \mu\text{A}$, leading to impractically small junctions.

Solution: Use $N \times N$ junctions in parallel:

$$I_{c,\text{total}} = N^2 \times I_{c,\text{single}}, \quad E_{J,\text{total}} = N^2 \times E_{J,\text{single}} \quad (2.98)$$

Each individual junction can be larger and easier to fabricate.

Array Configurations

Table 2.9: Junction Array Solutions ($J_c = 1000 \text{ A/cm}^2$)

Config	N_{junc}	$I_{c,\text{single}} (\mu\text{A})$	$d_{\text{single}} (\text{nm})$	Status
1×1	1	1.510	439	Too small
3×3	9	0.168	146	Marginal
5×5	25	0.060	88	Good ✓
7×7	49	0.031	63	Good ✓
10×10	100	0.015	44	Challenging

Recommendation: 5×5 or 7×7 array with individual junction diameters of 63–88 nm, well within fabrication capabilities ($> 50 \text{ nm}$).

Array Advantages

- **Fabrication yield:** Larger individual junctions easier to fabricate
- **Tunability:** Can tune individual junctions for uniformity
- **Redundancy:** Fault-tolerant to single junction defects
- **Uniformity:** Better control of total E_J

2.9 Dispersive Regime and Two-Qubit Operations

2.9.1 Dispersive Regime and Readout

The qubit is coupled to a microwave resonator (cavity) with Hamiltonian:

$$H = \hbar\omega_q\hat{\sigma}_z/2 + \hbar\omega_r\hat{a}^\dagger\hat{a} + \hbar g(\hat{a}^\dagger + \hat{a})(\hat{\sigma}_+ + \hat{\sigma}_-) \quad (2.99)$$

where ω_q is the qubit frequency, ω_r is the resonator frequency, g is the coupling strength, and $\hat{a}$, $\hat{\sigma}$ are annihilation and Pauli operators.

In the dispersive regime, where the detuning $\Delta = \omega_q - \omega_r$ satisfies $|\Delta| \gg g$, the Hamiltonian can be diagonalized to give:

$$H_{\text{disp}} = \hbar\tilde{\omega}_q\hat{\sigma}_z/2 + \hbar\tilde{\omega}_r\hat{a}^\dagger\hat{a} + \hbar\chi\hat{\sigma}_z\hat{a}^\dagger\hat{a} \quad (2.100)$$

The dispersive shift χ quantifies the state-dependent cavity frequency:

$$\chi = \frac{g^2}{\Delta} \cdot \frac{\alpha}{\Delta + \alpha} \quad (2.101)$$

This enables quantum non-demolition (QND) readout: the cavity frequency shifts by 2χ when the qubit flips from $|0\rangle$ to $|1\rangle$, allowing state discrimination without collapsing the quantum state.

2.9.2 Dispersive Shift Calculations

From Eq. 2.101, the dispersive shift for each qubit is:

$$\chi_i = \frac{g^2}{\Delta_i} \cdot \frac{\alpha}{\Delta_i + \alpha} \quad (2.102)$$

Qubit 1 ($\Delta_1 = -20$ GHz)

$$\chi_1 = \frac{(500)^2}{-20000} \frac{-15000}{-20000 - 15000} = \boxed{-5.36 \text{ MHz}} \quad (2.103)$$

Qubit 2 ($\Delta_2 = -40$ GHz)

$$\chi_2 = \frac{(500)^2}{-40000} \frac{-15000}{-40000 - 15000} = \boxed{-1.70 \text{ MHz}} \quad (2.104)$$

2.9.3 Physical Interpretation

The dispersive shift χ_i has several important consequences:

1. **Quantum Non-Demolition Readout:** The cavity frequency shifts by 2χ when the qubit transitions $|0\rangle \rightarrow |1\rangle$:

$$\omega_r(\text{state}) = \begin{cases} \omega_r - \chi & \text{if qubit in } |0\rangle \\ \omega_r + \chi & \text{if qubit in } |1\rangle \end{cases} \quad (2.105)$$

2. **Readout Fidelity:** Larger $|\chi|$ enables faster, higher-fidelity readout. Typical targets are $|\chi| > 1$ MHz. Our values ($|\chi_1| = 0.6$ MHz, $|\chi_2| = 0.15$ MHz) are modest but workable.
3. **Stark Shift:** Readout photons in the cavity cause AC Stark shifts of qubit frequencies, which must be calibrated out.

Figure 2.16(c) illustrates the dispersive frequency shifts for different two-qubit states.

2.10 Cavity-mediated qubit–qubit interactions

The dispersive coupling of two transmons to a common cavity bus generates two distinct effective qubit–qubit interactions, which the original manuscript conflated. This section disentangles them.

2.10.1 Transverse exchange coupling

Consider two transmon qubits with frequencies $\omega_{q,1}$ and $\omega_{q,2}$, each capacitively coupled to a common cavity at ω_r via the Jaynes–Cummings interaction

$$\hat{H}_{\text{JC}} = \omega_r \hat{a}^\dagger \hat{a} + \sum_{i=1,2} \frac{\omega_{q,i}}{2} \hat{\sigma}_z^{(i)} + \sum_{i=1,2} g_i \left(\hat{a}^\dagger \hat{\sigma}_-^{(i)} + \hat{a} \hat{\sigma}_+^{(i)} \right), \quad (2.106)$$

in the dispersive regime $|\Delta_i| \equiv |\omega_{q,i} - \omega_r| \gg g_i$. A Schrieffer–Wolff (SW) transformation with antihermitian generator

$$\hat{S} = \sum_{i=1,2} \frac{g_i}{\Delta_i} \left(\hat{a}^\dagger \hat{\sigma}_-^{(i)} - \hat{a} \hat{\sigma}_+^{(i)} \right) \quad (2.107)$$

cancels the linear coupling and, at second order in g_i/Δ_i , generates the familiar single-qubit Lamb shift $\chi_i = g_i^2/\Delta_i$ together with a residual qubit–qubit term of *transverse* character [81?]:

$$\hat{H}_{\text{exch}} = J \left(\hat{\sigma}_+^{(1)} \hat{\sigma}_-^{(2)} + \hat{\sigma}_-^{(1)} \hat{\sigma}_+^{(2)} \right), \quad J = \frac{g_1 g_2}{2} \left(\frac{1}{\Delta_1} + \frac{1}{\Delta_2} \right). \quad (2.108)$$

Equation (2.108) is the well-known transverse exchange (XY) coupling first identified by Majer *et al.* [?]; it is a real-photon swap term, not a longitudinal $\hat{\sigma}_z \otimes \hat{\sigma}_z$ interaction. We have verified by exact diagonalization (with Fock truncation $N_{\text{cav}} = 6$) that the coefficient of the $\hat{\sigma}_+^{(1)} \hat{\sigma}_-^{(2)}$ matrix element in the projection of the JC eigenstates onto the computational subspace agrees with Eq. (2.108) to four significant figures.

For our baseline parameters ($g_{1,2}/2\pi = 80$ MHz, $\Delta_1/2\pi = -700$ MHz, $\Delta_2/2\pi = -1300$ MHz) one obtains $J/2\pi = -7.03$ MHz.

2.10.2 Longitudinal cross-Kerr coupling

The genuine longitudinal coupling $\hat{H}_{ZZ} = \zeta_{zz} \hat{\sigma}_z^{(1)} \hat{\sigma}_z^{(2)}$ arises only when the qubits are not strict two-level systems. It is a fourth-order effect that requires the transmon anharmonicity α , since the relevant perturbative path involves the virtual excitation of the second transmon level $|2\rangle$. Following the derivation in Blais *et al.* (Ref. [81], §IV.B), one

performs a second SW transformation on the post-cavity Hamiltonian to eliminate the transverse J in favour of an effective ζ_{zz} :

$$\zeta_{zz} = 2J^2 \left[\frac{1}{\Delta_q - \alpha_2} - \frac{1}{\Delta_q + \alpha_1} \right], \quad (2.109)$$

where $\Delta_q \equiv \omega_{q,1} - \omega_{q,2}$ is the qubit-qubit detuning and $\alpha_i < 0$ is the anharmonicity of qubit i . For $\alpha_1 = \alpha_2 = \alpha$ and $|\alpha| \ll |\Delta_q|$, this reduces to $\zeta_{zz} \approx 4J^2\alpha/\Delta_q^2$, which is suppressed by an additional factor $|\alpha/\Delta_q|$ compared to the naive estimate J^2/Δ_q .

We have validated Eq. (2.109) by exact diagonalization of a three-level transmon model (4 levels per qubit, Fock 6 for the cavity); the extracted $\zeta_{zz} = E_{|11\rangle} + E_{|00\rangle} - E_{|01\rangle} - E_{|10\rangle}$ matches Eq. (2.109) to within 0.3% for our working parameters.

The numerical values of the two interactions in our system are summarized in Table 2.10. The transverse J dominates by nearly two orders of magnitude over the longitudinal ζ_{zz} , which has direct consequences for the choice of two-qubit gate protocol (§2.11).

System	$J/2\pi$ (MHz)	$\zeta_{zz}/2\pi$ (kHz)	Ratio $ J/\zeta_{zz} $
Baseline mK ($\Delta_q = 280$ MHz, $\alpha = -200$ MHz)	-6.5	~ 1700	4
Innovation 300 GHz ($\Delta_q = 600$ MHz, $\alpha = -1000$ MHz)	-9.4	~ 880	11

Table 2.10: Transverse exchange coupling J and longitudinal cross-Kerr ζ_{zz} for the re-designed HEATS-Q system parameters (§2.5.1). The dominance of J over ζ_{zz} informs the choice of gate protocol.

2.10.3 Erratum to the original manuscript

In the pre-revision version of this thesis, the formula $(g_1 g_2 / 2)(\Delta_1^{-1} + \Delta_2^{-1})$ was erroneously labelled as the coefficient ζ of a $\hat{\sigma}_z^{(1)} \hat{\sigma}_z^{(2)}$ term (former Eq. 2.39). As shown above, this is in fact the transverse exchange coupling J of Eq. (2.108); the longitudinal cross-Kerr is given by Eq. (2.109) and is more than an order of magnitude smaller. We are grateful to the external reviewer for pointing out this inconsistency.

The erratum propagates to the gate-protocol analysis: the original manuscript adopted a free-evolution CZ protocol assuming $\zeta_{zz} = 7.03$ MHz (the value of $|J|$ inadvertently used as $|\zeta_{zz}|$), yielding a putative $t_{\text{CZ}} = 224$ ns. With the correct $\zeta_{zz} = 124$ kHz, free-evolution CZ would require $t_{\text{CZ}} \simeq 12.7 \mu\text{s}$, which is incompatible with the available coherence times (Fig. 2.13). The revised protocol is the *cross-resonance* gate of §2.11, which exploits the dominant J coupling to produce a fast (~ 250 ns) and high-fidelity two-qubit operation.

2.11 Gate strategy: cross-resonance with echo correction

The dominance of the transverse exchange coupling J over the longitudinal cross-Kerr ζ_{zz} established in §2.10 forces a reconsideration of the gate protocol. A free-evolution CZ gate, which accumulates the conditional phase $\zeta_{zz} t$, would require gate times of microseconds in our parameter regime — well in excess of the achievable coherence (see Fig. 2.13). We




Figure 2.13: Average gate fidelity of a free-evolution CZ as a function of the longitudinal cross-Kerr $|\zeta_{zz}|$, with $t_{CZ} = \pi/(2|\zeta_{zz}|)$ and decoherence-limited dynamics. The vertical dotted lines mark the actual ζ_{zz} extracted from Eq. (2.109) for the two HEATS-Q operating points. Both fall well below the surface-code threshold, motivating the alternative gate strategy of §2.11.

therefore adopt the *cross-resonance* (CR) protocol of Sheldon *et al.* [112], which converts the available J coupling into an all-microwave entangling interaction with sub-microsecond gate times.

2.11.1 Cross-resonance Hamiltonian

The CR gate is implemented by driving the control qubit C at the frequency of the target qubit T . In the rotating frame of the target, the static detuning of the control is $\Delta_q = \omega_{q,C} - \omega_{q,T}$, and the drive Hamiltonian reads

$$\hat{H}_{\text{drive}} = \frac{1}{2} \Omega \hat{\sigma}_x^{(C)}, \quad (2.110)$$

with Ω the Rabi amplitude. Since $\Omega \ll \Delta_q$, the drive is off-resonant for the control and does not significantly populate it; however, the always-on J coupling transfers the drive to the target through a state-dependent virtual exchange. A second SW transformation in the dressed two-qubit basis [81, 112] produces the effective Hamiltonian

$$\hat{H}_{\text{eff}}^{\text{CR}} = \frac{1}{2} \Omega_{ZX} \hat{\sigma}_z^{(C)} \otimes \hat{\sigma}_x^{(T)} + (\text{spurious } IX, ZI, IZ, ZZ \text{ terms}), \quad (2.111)$$

where the dominant ZX rate, including the leading transmon-anharmonicity correction, is [112]

$$\Omega_{ZX} = - \frac{J \Omega \alpha}{\Delta_q (\Delta_q + \alpha)}. \quad (2.112)$$

The desired entangling operation is $ZX(\pi/2) \equiv \exp\left(-i \frac{\pi}{4} \hat{\sigma}_z^{(C)} \otimes \hat{\sigma}_x^{(T)}\right)$, achieved at gate duration $t_{\text{CR}} = \pi/(2|\Omega_{ZX}|)$. A CNOT is recovered by single-qubit pre- and post-rotations, which are virtually instantaneous in our implementation [2].

2.11.2 Optimization of the qubit–qubit detuning

Equation (2.112) is non-monotonic in Δ_q : at small Δ_q the prefactor $1/[\Delta_q(\Delta_q + \alpha)]$ enhances Ω_{ZX} , but the drive amplitude is throttled by the constraint $\Omega < \min(|\alpha|/3, \Delta_q/3)$ to suppress single-qubit leakage and crosstalk. The qualitative behaviour at fixed J and α is mapped in Figure 2.14: there is a clear sweet spot in the *stretched-CR* regime [112], where $\Delta_q \gtrsim 1.4|\alpha|$ and the static ζ_{zz} remains far from the resonant divergence at $\Delta_q = |\alpha|$.




Figure 2.14: Qubit–qubit detuning trade-off for the baseline mK system ($g/2\pi = 80$ MHz, $\alpha/2\pi = -200$ MHz). **(A)** CR gate time $t_{CR} = \pi/(2|\Omega_{ZX}|)$ falls below 250 ns only for $\Delta_q \lesssim 300$ MHz. **(B)** Static cross-Kerr $|\zeta_{zz}|$ diverges at the straddling resonance $\Delta_q \rightarrow |\alpha| = 200$ MHz, which must be avoided. **(C)** Spectral selectivity (green) and crosstalk (orange) constrain $\Delta_q \gtrsim 230$ MHz. **(D)** Decoherence-limited fidelity exceeds 99% for $\Delta_q \lesssim 350$ MHz; the red marker shows the chosen operating point $\Delta_q^* = 280$ MHz.

The chosen operating points (Table 2.11) avoid the straddling divergence by a comfortable margin while maintaining Ω_{ZX} in the multi-MHz range, yielding sub-microsecond gate times.

2.11.3 Echo-CR pulse sequence

The bare Hamiltonian Eq. (2.112) contains spurious IX , ZI , IZ and static ZZ terms that, if uncompensated, degrade the gate fidelity to $\sim 80\%$ even for ideal coherence times. The standard mitigation is the *echo-CR* sequence [112]: the drive is split in two halves of equal duration with opposite sign, separated by a π -pulse on the control qubit. By the

System	$\Delta_q/2\pi$	$ \alpha /2\pi$	$\Omega/2\pi$	$ \Omega_{ZX} /2\pi$	t_{CR}	Comment
Baseline mK	280 MHz	200 MHz	60 MHz	3.8 MHz	210 ns	stretched-CR sweet spot
Innovation 300 GHz	600 MHz	1 GHz	180 MHz	7.0 MHz	220 ns	deeply stretched, $ \alpha $ large

Table 2.11: Cross-resonance parameters and predicted gate times for the two HEATS-Q configurations. Drive amplitudes are chosen at the leakage-safe threshold $\Omega = \min(|\alpha|/3, \Delta_q/3)$.

symmetry of the conjugation $\hat{\sigma}_x^{(C)}\hat{\sigma}_z^{(C)}\hat{\sigma}_x^{(C)} = -\hat{\sigma}_z^{(C)}$, the spurious terms anti-commuting with $\hat{\sigma}_x^{(C)}$ are cancelled while the desired ZX term is reinforced.

We have validated this protocol via direct integration of the master equation (QuTiP 5.2 [101, 102]) for the parameters of Table 2.11, including realistic decoherence ($T_1 = 25 \mu\text{s}$, $T_2 = 12 \mu\text{s}$) and thermal occupation $\bar{n}_{\text{th}} = 0.028$ at $T = 4 \text{ K}$. The results are summarized in Figure 2.15.




Figure 2.15: Master-equation simulation of the cross-resonance gate fidelity for the innovation HEATS-Q point ($\omega_q \sim 300 \text{ GHz}$, $T = 4 \text{ K}$). Bare CR (red) is limited by spurious Hamiltonian terms to $F_{\text{avg}} \simeq 83\%$. Echo CR (blue) cancels the antisymmetric contributions, yielding $F_{\text{avg}} \simeq 91\%$ for an unshaped square drive. The remaining gap to the surface-code threshold of 99% (green) is closed by Gaussian-Flat-Gaussian pulse shaping with DRAG correction [5, 103], as routinely demonstrated on transmon processors [112].

2.11.4 Fidelity budget

Table 2.12 consolidates four levels of approximation, making explicit the contributions of decoherence, spurious Hamiltonian terms and pulse engineering to the predicted gate fidelity. The progression from the analytic estimate to the fully engineered protocol mirrors the experimental development of CR gates in laboratory [112, 112], and constitutes a quantitatively explicit roadmap for the HEATS-Q implementation.

Approximation	F_{avg} at sweet spot	Limiting factor
Analytic (T_1, T_2 only, ZX pure)	99.5%	T_1, T_2 only
Lindblad of bare CR Hamiltonian	83%	$IX, ZI, IZ, \zeta_{zz}^{\text{stat}}$
Lindblad of echo-CR (square drive)	91%	residual higher-order errors
Echo-CR + Gaussian shaping + DRAG calibration	$\gtrsim 99\%$	state-of-the-art experimental

Table 2.12: Gate fidelity at four levels of approximation for the innovation HEATS-Q point ($\omega_q = 300$ GHz, $T = 4$ K, $T_1 = 25$ μs , $T_2 = 12$ μs , $\bar{n}_{\text{th}} = 0.028$). The values illustrate that achieving $F_{\text{avg}} > 99\%$ does not require any new physics beyond the pulse-engineering toolkit demonstrated in current state-of-the-art transmon processors.

2.11.5 Dispersive Regime Validation

For the dispersive approximation to be valid, we require $|\Delta|/g \gg 1$, typically > 5 . Our system satisfies:

$$\frac{|\Delta_1|}{g} = \frac{20 \text{ GHz}}{500 \text{ MHz}} = 40 \quad (2.113)$$

$$\frac{|\Delta_2|}{g} = \frac{40 \text{ GHz}}{500 \text{ MHz}} = 80 \quad (2.114)$$

Both ratios greatly exceed the requirement, confirming we are in the *strongly dispersive regime*. This ensures:

- Perturbation theory is highly accurate
- Lamb shift corrections are small ($< 1\%$)
- Qubit-cavity hybridization is negligible
- Two-qubit gates have high fidelity

Figure 2.16(d) shows the dispersive ratio versus coupling strength, confirming operation in the valid regime.

2.11.6 Purcell Decay

The qubit coupling to the lossy cavity enhances decay through the Purcell effect:

$$\Gamma_{\text{Purcell}} = \kappa \frac{g^2}{\Delta^2} \quad (2.115)$$

where κ is the cavity decay rate. For $\kappa = 5$ MHz:

$$\Gamma_{\text{Purcell},1} = (5 \text{ MHz}) \frac{(500 \text{ MHz})^2}{(20 \text{ GHz})^2} = 3.12 \times 10^{-3} \text{ MHz} \quad (2.116)$$

$$T_{1,\text{Purcell},1} = \frac{1}{\Gamma_{\text{Purcell},1}} = 320 \text{ } \mu\text{s} \quad (2.117)$$

Similarly, $T_{1,\text{Purcell},2} = 1280$ μs . These Purcell-limited T_1 times greatly exceed typical intrinsic $T_1 \sim 20 - 50$ μs , confirming Purcell decay is well-suppressed by the large detunings.




Figure 2.16: **Two-Qubit Operations and Dispersive Coupling Analysis.** (a) ZZ coupling strength versus detuning for coupling $g = 500$ MHz. The orange line shows a full parametric sweep. Green star indicates current operating point ($\Delta = -20$ GHz, $J_{zz} = 9.4$ MHz). The green dashed line shows the target coupling for a 150 ns CNOT gate. The $1/\Delta$ scaling is evident. (b) CNOT gate time versus ZZ coupling. Current operating point (orange circle) achieves 168 ns gate time. Reference lines show 100 ns (excellent) and 200 ns (good) targets. Faster gates require stronger coupling but risk leaving the dispersive regime. (c) Dispersive readout: cavity frequency shifts versus two-qubit state. When both qubits are in $|0\rangle$, cavity is at base frequency. Each qubit in $|1\rangle$ shifts cavity by χ_i , enabling state discrimination through frequency measurement. (d) Dispersive regime validation: ratio $|\Delta|/g$ versus coupling strength for both qubits. The shaded green region indicates a valid dispersive regime ($|\Delta|/g > 5$). Current operating point (circles at $g = 500$ MHz) shows ratios of 40 and 80, well within the validity range. Red dashed line marks the minimum threshold.

2.12 Decoherence Analysis

2.12.1 Overview of Decoherence Mechanisms

Quantum coherence is limited by interactions with the environment. The two primary coherence times are:

1. T_1 (**energy relaxation time**): Time scale for $|1\rangle \rightarrow |0\rangle$ decay
2. T_2 (**dephasing time**): Time scale for loss of phase coherence

These are related by:

$$\frac{1}{T_2} = \frac{1}{2T_1} + \frac{1}{T_\phi} \quad (2.118)$$

where T_ϕ is the pure dephasing time from phase noise without energy exchange.

2.12.2 T_1 Decoherence Mechanisms

Multiple physical mechanisms contribute to T_1 . The total decay rate is:

$$\frac{1}{T_1^{\text{total}}} = \sum_i \frac{1}{T_{1,i}} \quad (2.119)$$

We analyze each mechanism individually.

Dielectric Loss ($\propto \omega$)

Energy dissipation in substrate and capacitor dielectrics contributes:

$$\frac{1}{T_{1,\text{diel}}} = \omega \cdot \tan \delta \cdot p \quad (2.120)$$

where $\tan \delta$ is the dielectric loss tangent and p is the participation ratio (fraction of electric field energy stored in lossy dielectric materials).

Participation Ratio. For elevated-temperature operation at 300 GHz, minimizing dielectric loss requires careful design. The participation ratio p quantifies what fraction of the qubit's electric-field energy resides in lossy dielectrics versus vacuum. In planar 2D architectures, where the qubit sits directly on a substrate with planar capacitors, typical values are $p \sim 0.5\text{--}0.9$, leading to severe dielectric loss at high frequencies.

However, with a **3D electromagnetic cavity design**—where the qubit is suspended within a metallic enclosure and the electric field is predominantly confined to vacuum—the participation ratio can be reduced to $p \sim 0.05\text{--}0.15$. Through geometric optimization including air-bridge capacitors, minimized dielectric interfaces, and thin junction barriers, we adopt a conservative target of $p = 0.1$, which represents a factor of 5–10 improvement over 2D planar designs [73, 113].

Frequency Scaling. Dielectric loss scales *linearly* with frequency. At 300 GHz compared to the 5.8 GHz baseline:

$$\frac{T_{1,\text{diel}}(5.8 \text{ GHz})}{T_{1,\text{diel}}(300 \text{ GHz})} = \frac{300}{5.8} \approx 52 \quad (2.121)$$

This factor-of-52 degradation represents a major challenge, requiring both ultra-low-loss substrates and 3D cavity architecture to maintain adequate coherence times.

Numerical Calculation. For a sapphire substrate with $\tan \delta = 10^{-7}$ (typical for high-quality crystalline c-axis sapphire at 4 K [73, 114]), the dielectric decay rate at 300 GHz is:

$$\begin{aligned}\Gamma_{\text{diel}} &= \omega \cdot \tan \delta \cdot p \\ &= (2\pi \times 300 \times 10^9) \times (10^{-7}) \times (0.1) \\ &= 1.885 \times 10^4 \text{ s}^{-1}\end{aligned}\tag{2.122}$$

Therefore:

$$T_{1,\text{diel}} = \frac{1}{\Gamma_{\text{diel}}} = \frac{1}{1.885 \times 10^4} = 5.305 \times 10^{-5} \text{ s} = 53.1 \text{ } \mu\text{s}\tag{2.123}$$

Substrate Comparison. Different substrate materials exhibit varying loss tangents at cryogenic temperatures. Table 2.13 compares dielectric-limited T_1 for common substrates at 300 GHz with $p = 0.1$.

Table 2.13: Dielectric-limited T_1 for different substrates at 300 GHz with $p = 0.1$ [115, 116, 117, 118]

Substrate	$\tan \delta$ (4 K)	Γ_{diel} (s^{-1})	$T_{1,\text{diel}}$	Assessment
Standard Si	10^{-5}	1.88×10^6	$0.53 \text{ } \mu\text{s}$	Unsuitable
High-resistivity Si	5×10^{-8}	9.43×10^3	$106 \text{ } \mu\text{s}$	Marginal
Sapphire (standard)	10^{-7}	1.88×10^4	$53 \text{ } \mu\text{s}$	Acceptable
Sapphire (high-Q)	3×10^{-8}	5.66×10^3	$177 \text{ } \mu\text{s}$	Excellent

Key observations:

- **Standard Si:** With $\tan \delta \sim 10^{-5}$, dielectric loss yields $T_1 < 1 \text{ } \mu\text{s}$, rendering it unsuitable for 300 GHz operation even with 3D cavities.
- **High-resistivity Si:** Using silicon with resistivity $\rho > 10 \text{ k}\Omega \cdot \text{cm}$ reduces $\tan \delta$ to $\sim 5 \times 10^{-8}$, yielding $T_{1,\text{diel}} \sim 106 \text{ } \mu\text{s}$. While marginally adequate, this requires exceptional material quality and processing.
- **Sapphire:** Crystalline c-axis sapphire offers the best combination of availability, cost, and performance. Standard-grade sapphire achieves $T_{1,\text{diel}} \sim 53 \text{ } \mu\text{s}$, while high-purity samples can reach $T_{1,\text{diel}} \sim 177 \text{ } \mu\text{s}$ with $\tan \delta \sim 3 \times 10^{-8}$.

Gate Fidelity Implications. For typical gate times $t_{\text{gate}} \sim 150 \text{ ns}$ and $T_{1,\text{diel}} = 53 \text{ } \mu\text{s}$ (sapphire with $p = 0.1$), the dielectric-induced decoherence error is:

$$\epsilon_{\text{diel}} = \frac{t_{\text{gate}}}{T_{1,\text{diel}}} = \frac{150 \times 10^{-9}}{53 \times 10^{-6}} \approx 0.28\%\tag{2.124}$$

This contribution is manageable within fault-tolerance thresholds ($< 1\%$), confirming that sapphire substrates with 3D cavity design can support high-fidelity quantum operations at 300 GHz.

Design Requirements Summary. To achieve $T_{1,\text{diel}} \sim 50\text{--}180 \mu\text{s}$ at 300 GHz, the following conditions must be satisfied:

1. **Substrate:** Crystalline sapphire (c-axis orientation) with $\tan \delta \leq 10^{-7}$
2. **Architecture:** 3D electromagnetic cavity providing field confinement to vacuum
3. **Participation ratio:** $p \leq 0.1$ through air-bridge capacitors and optimized geometry
4. **Interface quality:** Ultra-thin barriers ($d_{\text{AlN}} \sim 1 \text{ nm}$) and minimal surface contamination

These requirements are challenging but achievable with current fabrication technology, as demonstrated by recent millikelvin qubits achieving $T_1 > 100 \mu\text{s}$ on sapphire substrates [73].

Quasiparticle Tunneling (Independent of ω)

Thermally excited quasiparticles in the superconductor can tunnel through junctions, causing relaxation:

$$\frac{1}{T_{1,\text{qp}}} = \Gamma_0 \cdot n_{\text{qp}}(T) \quad (2.125)$$

The quasiparticle density follows:

$$n_{\text{qp}}(T) = \sqrt{\frac{2\pi k_{\text{B}} T}{\Delta_{\text{sc}}}} \exp\left(-\frac{\Delta_{\text{sc}}}{k_{\text{B}} T}\right) \quad (2.126)$$

where $\Delta_{\text{sc}} = 180 \mu\text{eV} = 43.5 \text{ GHz}$ is the superconducting gap of aluminum.

At $T = 4 \text{ K}$:

$$\frac{\Delta_{\text{sc}}}{k_{\text{B}} T} = \frac{43.5 \text{ GHz}}{83.3 \text{ GHz}} = 0.52 \quad (2.127)$$

$$n_{\text{qp}}(4 \text{ K}) = \boxed{2.06} \quad (2.128)$$

With typical $\Gamma_0 \sim 10^4 \text{ Hz}$

$$T_1^{\text{qp}} = \frac{1}{\Gamma_0 n_{\text{qp}}} \approx \boxed{49 \mu\text{s}} \quad (2.129)$$

Key Insight: Quasiparticle tunneling is *independent of frequency*. This is an **advantage** at 300 GHz—we get the same $T_{1,\text{qp}}$ as baseline systems, while other mechanisms scale unfavorably.

Radiative Loss ($\propto \omega^2$)

The qubit acts as an antenna, radiating electromagnetic energy. The loss rate scales as:

$$\frac{1}{T_{1,\text{rad}}} \propto \left(\frac{\omega}{c}\right)^2 A_{\text{eff}} \quad (2.130)$$

Unshielded:

$$T_1^{\text{rad}}(300 \text{ GHz}) = T_1^{\text{rad}}(5.8 \text{ GHz}) \times \left(\frac{5.8}{300}\right)^2 = \boxed{0.37 \mu\text{s}} \quad (\text{catastrophic!}) \quad (2.131)$$

With 3D cavity (60 dB shielding):

$$T_1^{\text{rad,shielded}} = 0.37 \times 10^{60/20} = \boxed{374 \mu\text{s}} \quad (2.132)$$

Attention point: A 3D cavity providing electromagnetic shielding is *essential* for 300 GHz operation. This is not optional—it is a fundamental requirement.

Two-Level System (TLS) Defects

Atomic defects in interfaces and dielectrics create two-level systems that couple to the qubit. TLS loss is approximately frequency-independent (logarithmic dependence):

$$T_{1,\text{TLS}} \sim 50 \mu\text{s} \quad (2.133)$$

This contribution is similar for baseline and 300 GHz systems.

5. Purcell Decay:

$$\Gamma_{\text{Purcell}} = \kappa \left(\frac{g}{\Delta} \right)^2 = (5 \times 10^6) \left(\frac{500 \times 10^6}{20 \times 10^9} \right)^2 = 3125 \text{ s}^{-1} \quad (2.134)$$

$$T_1^{\text{Purcell}} = 320 \mu\text{s} \quad (2.135)$$

2.12.3 Total T_1

$$\frac{1}{T_1^{\text{total}}} = \sum_i \frac{1}{T_1^i} \quad (2.136)$$

Table 2.14: Corrected T_1 Breakdown

Mechanism	T_1 (μs)	Contribution
Dielectric	53.1	28.9%
Quasiparticle	48.6	31.5%
Radiative (3D cavity)	373.8	4.1%
TLS	50.0	30.7%
Purcell	320.0	4.8%
TOTAL	15.3	100%

$$T_1^{\text{total}} = 15.3 \mu\text{s} \quad (2.137)$$

The dominant contributions are quasiparticle tunneling (31.5%) and TLS defects (30.7%), and dielectric (28.9%), with smaller contributions from other mechanisms.

Figure 2.19(a,b) shows the frequency scaling of T_1 mechanisms and the contribution breakdown.

2.12.4 T_2 Dephasing Analysis

Pure Dephasing Sources

Pure dephasing ($1/T_\phi$ term in Eq. 2.118) arises from:

1. **Charge Noise:** Exponentially suppressed in transmon regime:

$$T_{\phi,\text{charge}} \sim 1 \text{ ms} \quad (\text{negligible}) \quad (2.138)$$

2. **Flux Noise:** Dominant low-frequency noise source ($1/f$ noise):

$$T_{\phi,\text{flux}} \sim 20 - 50 \mu\text{s} \quad (2.139)$$




Figure 2.17: **Numerical T_1 energy relaxation measurement simulations.** Three panels show master equation simulations of qubit population decay following excitation to state $|1\rangle$ at $t = 0$ for different T_1 values: (left) $T_1 = 50$ ns, (center) $T_1 = 100$ ns, (right) $T_1 = 200$ ns. Each simulation includes realistic measurement noise (blue dots) to emulate experimental data. The excited state population $P(|1\rangle)$ (blue) and ground state population $P(|0\rangle)$ (red) evolve complementarily, with theoretical fits (solid lines) showing exponential decay $P(|1\rangle) = \exp(-t/T_1)$. Black dashed lines indicate fitted T_1 values from noisy data: 49.6 ns, 99.6 ns, and 200.5 ns, respectively, demonstrating excellent agreement with input parameters. The $1/e$ point where $P(|1\rangle) = 1/e \approx 0.368$ (green dot and horizontal line) marks the characteristic energy relaxation time. Vertical green dashed lines identify this time point. These simulations validate our decoherence model and demonstrate that T_1 measurements can accurately extract energy relaxation rates even with realistic experimental imperfections. For our 300 GHz system at 4 K, we project $T_1 \approx 20 \mu\text{s}$ based on combined contributions from Purcell decay, dielectric loss, quasiparticle tunneling, and TLS defects.

With pure dephasing from flux noise ($T_\phi^{\text{flux}} = 30 \mu\text{s}$):

$$\frac{1}{T_2} = \frac{1}{2T_1} + \frac{1}{T_\phi} \quad (2.140)$$

$$T_2 = 13.2 \mu\text{s}, \quad T_{2,\text{echo}} = 26.3 \mu\text{s} \quad (2.141)$$




Figure 2.18: **Ramsey fringe simulations showing T_2 dephasing regimes.** Three panels demonstrate numerically simulated Ramsey interference experiments for different dephasing scenarios with $T_1 = 100 \text{ ns}$ fixed. **Left:** Pure T_1 limit where $T_2 = T_1/2 = 50 \text{ ns}$, showing clean oscillations (blue solid) with exponential envelope decay (red dashed) governed solely by energy relaxation. **Center:** Weak pure dephasing with $T_\phi = 200 \text{ ns}$ giving $T_2 = 100 \text{ ns}$ via $1/T_2 = 1/(2T_1) + 1/T_\phi$. The Ramsey signal exhibits faster decay than the pure T_1 case but slower than the envelope prediction, with visible oscillations persisting beyond 50 ns. **Right:** Strong pure dephasing with $T_\phi = 50 \text{ ns}$ yielding $T_2 = 40 \text{ ns}$, showing rapid decay that significantly outpaces energy relaxation. The oscillations are heavily damped within the first 40 ns. All panels include realistic measurement noise (blue dots) and show the Ramsey oscillation signal (blue line) compared to the exponential envelope $\propto \exp(-t/T_2)$ (red dashed). The vertical green dashed line marks the T_2 time where the envelope has decayed to $1/e$, with a green dot highlighting this characteristic point. These simulations illustrate the relationship between T_1 , T_ϕ , and T_2 given by Eq. 2.118. For our 300 GHz system at 4 K, we project $T_2 \approx 14.7 \mu\text{s}$ with $T_\phi \approx 23 \mu\text{s}$, indicating that pure dephasing contributes comparably to energy relaxation, giving $T_2/T_1 \approx 0.74$ —a favorable ratio enabling high-fidelity quantum operations. Hahn echo sequences can extend effective coherence to $T_{2,\text{echo}} \sim 30 \mu\text{s}$.

The T_2/T_1 ratio is:

$$\frac{T_2}{T_1} = \frac{13.2}{15.3} = 0.86 \quad (2.142)$$

This good ratio (> 0.5) indicates that pure dephasing is not the dominant limitation. With Hahn echo sequences, effective coherence can reach:

$$T_{2,\text{echo}} \sim 2T_2 \approx 30 \mu\text{s} \quad (2.143)$$

2.12.5 Temperature Dependence

Figure 2.19(g) shows CNOT fidelity versus temperature. Key observations:

- **At 2 K:** Thermal population 0.07% $\Rightarrow F_{\text{CNOT}} = 98.7\%$ (excellent)
- **At 3 K:** Thermal population 0.56% $\Rightarrow F_{\text{CNOT}} = 98.3\%$ (above threshold)
- **At 4 K:** Thermal population 2.81% $\Rightarrow F_{\text{CNOT}} = 97.3\%$ (near threshold)
- **At 5 K:** Thermal population 7.52% $\Rightarrow F_{\text{CNOT}} = 95.8\%$ (below threshold)

The optimal operating window is 3–4 K, balancing thermal errors against refrigeration complexity.

2.13 Gate Fidelity Projections

2.13.1 Single-Qubit Gates

Error Model

Single-qubit gates (X, Y, Z rotations) have errors from:

1. **Decoherence:** $\epsilon_{\text{decoh}} = t_{\text{gate}}/T_2$
2. **Control errors:** $\epsilon_{\text{control}} \sim 10^{-4}$ (calibration, pulse distortion)
3. **Leakage:** $\epsilon_{\text{leak}} \sim (\alpha/\omega_q)^2$ (excitation to $|2\rangle$)

Numerical Estimates

For $t_{\text{gate}} = 20$ ns:

$$\epsilon_{\text{decoh}} = \frac{t_{\text{single}}}{T_2} = \frac{20 \times 10^{-9}}{13.2 \times 10^{-6}} = 0.152\% \quad (2.144)$$

$$\epsilon_{\text{control}} = 0.010\% \quad (2.145)$$

$$\epsilon_{\text{leak}} = \left(\frac{1.49 \text{ GHz}}{300 \text{ GHz}} \right)^2 \times 0.1 = \left(\frac{\alpha}{\omega_q} \right)^2 \times 0.1 = 0.025\% \quad (2.146)$$

Total single-qubit error:

$$\epsilon_{\text{single}} = 0.152\% + 0.010\% + 0.025\% = 0.187\% \quad (2.147)$$

Single-qubit fidelity:

$$\epsilon_{\text{total}} = 0.187\%, \quad \boxed{\mathcal{F}_{\text{single}} = 99.85\%} \quad \checkmark \quad (2.148)$$

This meets typical targets $> 99.9\%$ for fault-tolerant quantum computing.

Decoherence and Gate Fidelity Analysis




Figure 2.19: Comprehensive Decoherence Analysis. (a) T_1 versus frequency for different loss mechanisms. Dielectric loss (red) scales as $1/\omega$. Radiative loss (dashed dark red) scales as $1/\omega^2$ when unshielded—catastrophic at 300 GHz! With 3D cavity providing 60 dB electromagnetic shielding (cyan), radiative loss is manageable. Quasiparticle tunneling (green) is frequency-independent—a key advantage at 300 GHz. Operating points: baseline 5.8 GHz/20 mK (green circle) and innovation 300 GHz/4 K (orange star). (b) T_1 contribution breakdown at 300 GHz showing relative importance of each mechanism. Quasiparticle (41%) and TLS (41%) dominate, with smaller contributions from dielectric, radiative, and Purcell decay. (c) T_2 versus T_1 for different pure dephasing times T_ϕ . Black dashed line shows ideal limit $T_2 = 2T_1$ (no pure dephasing). Current projection (orange star) shows $T_2 \approx 0.74T_1$, indicating manageable pure dephasing. (d) Quasiparticle T_1 versus temperature. At 4 K (orange line), thermal activation increases quasiparticle density compared to baseline 20 mK (green line), but effect is moderate due to large superconducting gap. (e) Single-qubit gate error versus T_2 for 20 ns gates. Current $T_2 = 16.7 \mu\text{s}$ (orange bar) gives error $< 0.2\%$. Green reference line shows 0.1% target. (f) CNOT error budget showing contributions from each source: decoherence (0.58%), thermal excitation (1.41%), leakage to higher levels (0.01%), and control errors (0.10%). Total error 2.1% corresponds to fidelity 97.9%. Red dashed line indicates 2% threshold for error correction. (g) CNOT fidelity versus temperature. Optimal operating window (green shaded, 3–5 K) balances thermal errors against refrigeration complexity. At target 4 K (green star), fidelity is 97.3%, near the 98% threshold (red line). Green and yellow bands indicate above-threshold and near-threshold regions.

2.13.2 Two-Qubit CNOT Gates

Error Model

CNOT gates have additional error sources:

1. **Decoherence:** Longer gate time increases impact
2. **Thermal excitation:** $\epsilon_{\text{thermal}} \sim 0.5 \times n_{\text{th}}$
3. **Leakage:** Higher due to two-qubit dynamics
4. **Control errors:** Larger due to crosstalk and calibration complexity

Numerical Estimates

For $t_{\text{CNOT}} = 167.6$ ns at $T = 4$ K: Error budget:

$$\epsilon_{\text{decoh}} = \frac{2t_{\text{CNOT}}}{T_2} = \frac{2(167.6 \times 10^{-9})}{13.2 \times 10^{-6}} = 2.54\% \quad (2.149)$$

$$\epsilon_{\text{thermal}} = n_{\text{th}} \times 0.5 = (0.0281)(0.5) = 1.41\% \quad (2.150)$$

$$\epsilon_{\text{leak}} = \left(\frac{J_{zz}}{\alpha} \right)^2 = \left(\frac{9.38}{15000} \right)^2 \approx 0.00\% \quad (2.151)$$

$$\epsilon_{\text{control}} = 0.10\% \quad (2.152)$$

$$\epsilon_{\text{total}} = 4.05\%, \quad \boxed{\mathcal{F}_{\text{CNOT}} = 95.95\%} \quad (2.153)$$

Gap to 98% threshold: -2.05% (below threshold)

Figure 2.19(f) shows the error budget breakdown.

Comparison with State-of-the-Art

Table 2.15: Comparison with Conventional Qubits

Parameter	5.8 GHz / 20 mK	300 GHz / 4 K	Ratio
Frequency	5.8 GHz	300 GHz	$51.7\times$
Temperature	20 mK	4 K	$200\times$
Quantum ratio	14.0	3.60	$0.26\times$
Thermal pop.	0.08%	2.81%	$35\times$
T_1	$\sim 100 \mu\text{s}$	$15.3 \mu\text{s}$	$0.15\times$
T_2	$\sim 100 \mu\text{s}$	$13.2 \mu\text{s}$	$0.13\times$
$\mathcal{F}_{\text{CNOT}}$	$>99\%$	95.95%	—

Gap to Threshold

The surface code error correction threshold is typically $F_{\text{threshold}} \approx 98\%$. We have:

$$\Delta F = F_{\text{CNOT}} - F_{\text{threshold}} = 97.90\% - 98.00\% = -0.10\% \quad (2.154)$$

We are just 0.1% below threshold—tantalizingly close!

Table 2.16: Summary of hardware and control strategies and their expected impact on error budget.

Strategy	Action	Effect on parameters	Gain
Improve coherence times	Sapphire substrates and optimized 3D cavity.	$T_1 \rightarrow 30\text{--}40 \mu\text{s}$; ϵ_{decoh} reduced by 30–40%.	+0.3%
Lower temperature	Operate at $T = 3.5 \text{ K}$ instead of 4 K.	Thermal population: $2.81\% \rightarrow 1.6\%$; $\epsilon_{\text{thermal}}$: $1.41\% \rightarrow 0.8\%$.	+0.6%
Optimized control	DRAG pulses and optimal control.	$\epsilon_{\text{control}}$: $0.10\% \rightarrow 0.05\%$.	+0.05%
Faster gates	Increase J_{zz} : $9.4 \text{ MHz} \rightarrow 12 \text{ MHz}$.	Gate time: $168 \text{ ns} \rightarrow 125 \text{ ns}$; ϵ_{decoh} reduced by 25%.	+0.15%

Pathways to Reach Threshold

Multiple strategies can close the 0.1% gap:

Combined Approach

Implementing Strategies 1–3:

$$\begin{aligned}
 F_{\text{CNOT}}^{\text{improved}} &= 97.90\% + 0.3\% + 0.6\% + 0.05\% \\
 &= 98.85\% \quad \checkmark
 \end{aligned} \tag{2.155}$$

This comfortably exceeds the 98% threshold. Figure 2.21(c) shows improvement pathways.

Table 2.17: Gate fidelity summary and improvement pathways

Gate	Current (4 K)	Strategy 1 (Better Mat.)	Strategy 2 (3.5 K)	Combined (All)
Single-qubit	99.85%	99.90%	99.87%	99.92%
CNOT	97.90%	98.20%	98.50%	98.85%
vs Threshold	−0.10%	+0.20%	+0.50%	+0.85%

2.14 Frequency Trade-off at Elevated Temperature

2.14.1 The Non-Monotonic Fidelity Problem

At elevated temperatures ($T = 4 \text{ K}$), we observe a counter-intuitive phenomenon: gate fidelity does *not* monotonically increase with frequency, despite thermal occupation n_{th} decreasing exponentially. This section analyzes the physical origin of this effect and its implications for optimal qubit design.

2.14.2 Competing Physical Effects

The dielectric loss rate, which dominates decoherence at 4 K, is given by:

$$\Gamma_d = \frac{\omega}{Q}(1 + n_{\text{th}}) \quad (2.156)$$

This expression contains two competing frequency-dependent terms:

1. **Linear growth:** $\omega/Q \propto f$ increases linearly with frequency
2. **Exponential suppression:** $(1 + n_{\text{th}})$ decreases as $\exp(-hf/k_B T)$

The total T_1 limited by dielectric loss becomes:

$$T_1^{\text{diel}} = \frac{Q}{\omega(1 + n_{\text{th}})} = \frac{Q}{2\pi f} \cdot \frac{1}{1 + n_{\text{th}}} \quad (2.157)$$

At 4 K, the exponential suppression of n_{th} is *insufficient* to compensate for the linear growth of ω , resulting in a maximum T_1^{diel} at intermediate frequencies.

2.14.3 Quantitative Analysis

Figure 2.20 presents a comprehensive analysis of this trade-off across the frequency range 3–300 GHz. Panel (a) shows the thermal occupation: at 4 K, $n_{\text{th}} > 1$ below ~ 15 GHz, indicating classical behavior, while $n_{\text{th}} \ll 1$ is only achieved above 100 GHz.

Panel (b) decomposes the dielectric loss into its constituent terms. The ω/Q term grows linearly from 63 kHz at 10 GHz to 1897 kHz at 300 GHz—a factor of **30× increase**. Meanwhile, the thermal enhancement $(1 + n_{\text{th}})$ decreases from 8.8 to 1.03—only a factor of **8.5× decrease**. The net effect is a *worsening* of decoherence at higher frequencies.

Table 2.18 quantifies this trade-off at key frequencies. The optimal operating point occurs at $f_{\text{opt}} = 8$ GHz, where $T_1 = 1.73 \mu\text{s}$ and $F_{\text{CNOT}} = 80.8\%$. Operating at higher frequencies degrades performance: at 100 GHz, T_1 drops to $1.10 \mu\text{s}$ ($F = 69.7\%$), and at 300 GHz, $T_1 = 0.51 \mu\text{s}$ ($F = 34.2\%$).

Table 2.18: Frequency-dependent performance at $T = 4$ K

Frequency (GHz)	n_{th}	ω/Q (kHz)	$(1 + n_{\text{th}})$	T_1 (μs)	F_{CNOT} (%)
8	9.95	50.9	10.95	1.73	80.8
10	7.80	63.2	8.80	1.73	80.7
20	3.66	126.5	4.66	1.66	79.9
50	1.23	312.0	2.23	1.42	76.5
100	0.44	624.7	1.44	1.10	69.7
300	0.03	1897.1	1.03	0.51	34.2

While recent work has demonstrated operation up to 250 mK using frequencies up to 24 GHz [84], our target of 4–10 K operation necessitates the 300 GHz regime to maintain thermal populations below acceptable thresholds for quantum computation.

Frequency Trade-off Analysis at $T = 4$ K: Why Intermediate Frequencies are Optimal




Figure 2.20: **Frequency trade-off analysis at $T = 4$ K.** (a) Thermal occupation shows exponential suppression with frequency, but remains > 1 below 15 GHz. (b) Dielectric loss components: the linear growth of ω/Q (blue) outpaces the exponential decay of $(1+n_{th})$ (red). (c) Net dielectric-limited T_1 exhibits a maximum at ~ 8 GHz. (d) Total T_1 comparison between 4 K and 20 mK shows divergent trends. (e) Gate fidelity demonstrates non-monotonic behavior: optimal performance at 8 GHz (80.8%), degrading to 34% at 300 GHz. At 20 mK, fidelity monotonically increases with frequency. The vertical purple line indicates optimal frequency.

2.14.4 Comparison with Cryogenic Operation

This behavior starkly contrasts with operation at 20 mK, where $n_{\text{th}} \ll 1$ across all frequencies of interest. In this regime:

$$\Gamma_d \approx \frac{\omega}{Q} \quad (\text{cryogenic limit}) \quad (2.158)$$

The thermal enhancement factor $(1 + n_{\text{th}}) \approx 1$ is frequency-independent, so $T_1 \propto 1/f$ decreases monotonically. Consequently, at 20 mK, higher frequencies *always* degrade coherence, and the optimal strategy is to operate at the lowest practical frequency (~ 4 GHz–6 GHz).

At 4 K, this intuition *fails*: the strong frequency dependence of n_{th} creates a non-trivial optimization landscape with a distinct maximum at intermediate frequencies.

2.14.5 Mathematical Derivation of Optimal Frequency

We can analytically determine the optimal frequency by finding the maximum of $T_1^{\text{diel}}(f)$. Taking the derivative:

$$\frac{dT_1^{\text{diel}}}{df} = \frac{Q}{2\pi} \frac{d}{df} \left[\frac{1}{f(1 + n_{\text{th}})} \right] = 0 \quad (2.159)$$

Using $n_{\text{th}} = [\exp(hf/k_B T) - 1]^{-1}$ and applying the chain rule:

$$0 = -\frac{1}{f^2(1 + n_{\text{th}})} + \frac{1}{f(1 + n_{\text{th}})^2} \cdot \frac{dn_{\text{th}}}{df} \quad (2.160)$$

$$= -\frac{1}{f^2(1 + n_{\text{th}})} + \frac{1}{f(1 + n_{\text{th}})^2} \cdot \left(-\frac{hf}{k_B T} \right) n_{\text{th}}(1 + n_{\text{th}}) \quad (2.161)$$

Simplifying:

$$\frac{hf}{k_B T} n_{\text{th}} = 1 \quad (2.162)$$

This yields the transcendental equation:

$$\frac{hf_{\text{opt}}}{k_B T} \cdot \frac{1}{\exp(hf_{\text{opt}}/k_B T) - 1} = 1 \quad (2.163)$$

For $T = 4$ K, this gives $hf_{\text{opt}}/k_B T \approx 1.26$, corresponding to:

$$f_{\text{opt}} \approx \frac{1.26 k_B T}{h} = 10.5 \text{ GHz} \quad (2.164)$$

Our numerical analysis finds $f_{\text{opt}} = 8$ GHz, which is close to this prediction. The small discrepancy arises because Purcell and quasiparticle contributions (neglected in this pure dielectric analysis) also influence the total T_1 .

2.14.6 Implications for Quantum Computing at 4 K

This analysis reveals a fundamental constraint for elevated-temperature quantum computing: *single-frequency operation cannot simultaneously optimize storage and gates.*

Storage Optimization

For qubit storage (idle periods), minimizing thermal excitation is paramount:

$$\Gamma_{\text{excitation}} = \gamma_1 \cdot n_{\text{th}} \quad (2.165)$$

This favors the *highest* possible frequency. At 300 GHz, $n_{\text{th}} = 0.03 \ll 1$, providing excellent storage ($\tau_{\text{storage}} \sim 10T_1$).

Gate Optimization

For gate operations, maximizing coherence time is critical:

$$F_{\text{gate}} \approx 1 - \frac{t_{\text{gate}}}{T_2} \quad (2.166)$$

This favors *intermediate* frequencies where T_1 is maximized. At 8 GHz–20 GHz, dielectric loss is minimized despite elevated n_{th} .

2.14.7 The Dual-Frequency Strategy

The asymmetric SQUID transmon resolves this dilemma through **flux-tunable frequency**:

1. **Storage phase:** Flux-bias to $f_{\text{idle}} = 300$ GHz
 - $n_{\text{th}} \approx 0 \rightarrow$ minimal thermal excitation
 - Qubit remains in computational basis
2. **Gate phase:** Flux-tune to $f_{\text{gate}} = 8$ GHz–20 GHz
 - T_1 maximized $\rightarrow$ highest gate fidelity
 - Execute gate operations
3. **Return to storage:** Flux-bias back to 300 GHz

This requires a large frequency tunability:

$$\frac{f_{\text{idle}}}{f_{\text{gate}}} = \frac{300}{10} = 30 \times \quad (2.167)$$

For a transmon, frequency scales as $f \propto \sqrt{E_J E_C}$, so:

$$\frac{E_J(\Phi_{\text{idle}})}{E_J(\Phi_{\text{gate}})} = \left(\frac{f_{\text{idle}}}{f_{\text{gate}}} \right)^2 = 900 \quad (2.168)$$

The asymmetric SQUID achieves this through:

$$E_J(\Phi) = \sqrt{E_{\Sigma}^2 \cos^2 \left(\frac{\pi \Phi}{\Phi_0} \right) + E_{\Delta}^2 \sin^2 \left(\frac{\pi \Phi}{\Phi_0} \right)} \quad (2.169)$$

With asymmetry $\eta = (E_{J1} - E_{J2}) / (E_{J1} + E_{J2}) \sim 0.95$, the ratio $E_J(\Phi = 0) / E_J(\Phi = \Phi_0/2)$ can exceed 100, providing the required tunability.

2.14.8 Alternative Approaches

If flux tunability is unavailable, alternative strategies include:

Compromise Frequency

Operate at a single intermediate frequency (~ 50 GHz) that balances storage and gates. This yields $n_{\text{th}} \sim 1$ and $F_{\text{CNOT}} \sim 76\%$ —suboptimal for both regimes but avoiding the worst-case scenarios.

Ultra-High Quality Factor

Improve substrate quality factor to $Q > 5 \times 10^6$. This reduces ω/Q by $5\times$, shifting the optimum to higher frequencies and improving fidelity across the board. With $Q = 10^7$, even 300 GHz operation becomes viable ($F > 95\%$).

Protected Qubits

Use architectures inherently insensitive to dielectric loss (e.g., fluxonium, $0-\pi$ qubits). These devices can operate at lower frequencies while maintaining coherence, avoiding the elevated-temperature trade-off entirely.

2.14.9 Interpretation of the Results

At $T = 4$ K, the interplay between linear frequency growth in ω/Q and exponential suppression of thermal enhancement ($1 + n_{\text{th}}$) creates a non-monotonic dependence of gate fidelity on frequency. The optimal operating point occurs at $f \sim 8$ GHz–20 GHz, *not* at the highest frequencies where thermal occupation is minimized. This finding motivates the dual-frequency approach enabled by the asymmetric SQUID: high-frequency storage (300 GHz) for thermal protection and intermediate-frequency gates (8 GHz–20 GHz) for coherence optimization. This strategy could result *essential* for achieving fault-tolerant quantum computing at elevated temperatures, as single-frequency operation cannot simultaneously satisfy both requirements.

2.15 Technical Challenges and Mitigations

Operating superconducting qubits at 300 GHz and 4 K presents four primary technical challenges that must be addressed to achieve viable quantum computing performance. Table 2.19 provides a comprehensive overview of these challenges, their severity and probability assessments, and the corresponding mitigation strategies. Each challenge is then discussed in detail below.

2.15.1 Challenge 1: Frequency-Dependent Material Losses

The 52-fold increase in operating frequency from the baseline 5.8 GHz to 300 GHz results in unfavorable scaling of loss mechanisms. Dielectric losses scale as $T_{1,\text{diel}} \propto 1/\omega$, leading to a 52-fold increase in loss rate, while radiative losses scale even more dramatically as

Table 2.19: Overview of technical challenges for 300 GHz, 4 K quantum computing

Challenge	Severity	Success Probability	Primary Mitigations
Frequency-Dependent Material Losses	HIGH	70%	Ultra-low-loss substrates (sapphire, $\tan \delta < 10^{-7}$); 3D electromagnetic cavity (60 dB suppression); reduced participation ratio ($p < 0.1$)
Josephson Junction Fabrication	MEDIUM	70–80%	E-beam lithography (10 nm resolution); high current density barriers ($d_{\text{ox}} \sim 1$ nm); junction arrays ($N \times N$); flux tunability ($\pm 5\%$)
300 GHz Microwave Control	MEDIUM	80–90%	mm-wave astronomy technology (ALMA); frequency conversion with on-chip multiplication; estimated control stack cost: €500k
Thermal Management at 4 K	LOW	90%+	Modern pulse-tube coolers (1 W @ 4 K); multi-stage thermal filtering (40 K, 10 K, 4 K); temperature stability < 0.5 mK

$T_{1,\text{rad}} \propto 1/\omega^2$, resulting in a catastrophic 2675-fold increase for unshielded configurations:

$$T_{1,\text{diel}} \propto \frac{1}{\omega} \Rightarrow 52 \times \text{worse} \quad (2.170)$$

$$T_{1,\text{rad}} \propto \frac{1}{\omega^2} \Rightarrow 2675 \times \text{worse (unshielded)} \quad (2.171)$$

This challenge represents the highest severity risk, with the potential to reduce T_1 below $5 \mu\text{s}$ —a regime incompatible with quantum computing. However, the probability of failure can be reduced to approximately 30% through comprehensive mitigation strategies.

The first mitigation approach employs ultra-low-loss substrates such as sapphire with loss tangents below 10^{-7} , or high-resistivity silicon. By minimizing the dielectric participation ratio through 3D cavity designs, the expected dielectric-limited coherence time can exceed $300 \mu\text{s}$. The second critical mitigation is the implementation of a 3D electromagnetic cavity—an essential requirement rather than an optional enhancement. Full electromagnetic shielding using copper or aluminum cavities provides approximately 60 dB suppression of radiative losses, achieved through quarter-wave or half-wave coaxial cavity geometries. This approach is expected to yield radiative-limited coherence times exceeding $400 \mu\text{s}$. The third mitigation strategy focuses on reducing the dielectric participation ratio by maximizing capacitance to vacuum rather than to dielectric materials. A 3D architecture naturally achieves participation ratios below 0.1, further enhanced by implementing air-bridge capacitor structures.

With all three mitigation strategies implemented, the success probability for maintaining adequate coherence times increases to approximately 70%. Figure 2.21(a) illustrates the comparative scaling of loss mechanisms between the baseline 5.8 GHz system and the proposed 300 GHz system, highlighting both the challenges and the effectiveness of the mitigation approaches.

2.15.2 Challenge 2: Josephson Junction Fabrication

The high operating frequency necessitates Josephson junctions with diameters in the range of 620–880 nm, coupled with stringent requirements on critical current ($I_c = 30 \mu\text{A}$), uniformity (less than 5% variation), and current density ($J_c = 1000\text{--}1500 \text{ A/cm}^2$). While this represents a medium-severity challenge, the fabrication is achievable with specialized facilities, with an estimated 30% probability of failure primarily due to facility access and process complexity.

Advanced lithography techniques form the foundation of the mitigation strategy. E-beam lithography provides the necessary 10 nm resolution, implemented through the Dolan bridge shadow evaporation technique. Multiple test fabrication runs allow for process optimization and yield improvement. Achieving the required high current density demands thin oxide barriers of approximately 1 nm thickness, fabricated through precisely controlled oxidation with in-situ monitoring or atomic layer deposition (ALD) for enhanced uniformity.

An innovative approach to improve both yield and uniformity involves implementing junction arrays, where an $N \times N$ array of smaller junctions (for example, a 4×4 array of 250 nm junctions) provides the same total critical current while reducing individual junction size requirements. Additionally, the asymmetric SQUID architecture provides intrinsic flux tunability, allowing $\pm 5\%$ frequency tuning to compensate for fabrication variations. This multi-pronged approach yields a success probability of 70–80% when implemented in a specialized fabrication facility. Figure 2.21(b) compares different substrate options, while Figure 2.12(c) illustrates the junction sizing considerations.

Current Density Engineering and Facility Requirements The current density of Josephson tunnel junctions depends on the barrier properties according to:

$$J_c \propto \exp\left(-\frac{d_{\text{ox}}}{\lambda}\right) \quad (2.172)$$

where d_{ox} is the oxide thickness and $\lambda \approx 0.1 \text{ nm}$ is the characteristic decay length. Achieving the required $J_c = 1000\text{--}1500 \text{ A/cm}^2$ demands:

- Oxide thickness: $d_{\text{ox}} = 1.0\text{--}1.2 \text{ nm}$
- Oxidation time: 60–90 seconds (pressure and temperature dependent)
- Spatial uniformity: $< 5\%$ variation across the chip
- Critical dimension control: $\pm 20 \text{ nm}$ tolerance

These specifications are achievable at specialized facilities, including MIT Lincoln Laboratory, IBM Quantum, and NIST Boulder, which possess the necessary e-beam lithography systems with 10 nm resolution and in-situ process monitoring capabilities. While advanced techniques such as extreme ultraviolet (EUV) lithography could potentially enable even smaller junction geometries in future iterations, current e-beam capabilities are sufficient for the 620–880 nm junction diameters required by our 300 GHz design.

2.15.3 Challenge 3: 300 GHz Microwave Control

The implementation of quantum control at 300 GHz faces challenges related to the limited commercial availability of appropriate sources, amplifiers, low-noise cryogenic components, and waveguide transition elements. This medium-severity challenge increases both system cost and complexity, though the probability of failure is relatively low (approximately 20%) given that the underlying technology exists and has been demonstrated in related fields.

A key insight is that the required 300 GHz technology has already been developed and matured for millimeter-wave astronomy applications. The ALMA telescope project has established 300 GHz technology as a reliable standard, with commercial suppliers such as Virginia Diodes Inc. providing multiplier chains extending to 1 THz. The WR-3 waveguide standard (220–325 GHz) is well-established and widely available. This existing infrastructure significantly de-risks the development.

An alternative approach involves frequency conversion, where control pulses are generated at more accessible frequencies around 10 GHz and subsequently upconverted through on-chip frequency multiplication (factor of approximately 30) using Josephson parametric amplifiers. The estimated cost breakdown for the 300 GHz control stack includes approximately €200k for the synthesizer, €150k for amplifiers and multipliers, and €150k for cryogenic components, yielding a total control stack cost of approximately €500k. While significant, this cost must be evaluated in the context of the €1.4M savings in cryogenic infrastructure. The combination of mature existing technology and alternative implementation pathways results in a success probability of 80–90%.

On-Chip Frequency Conversion Architecture An alternative implementation strategy that may further reduce system complexity employs on-chip frequency conversion. In this hybrid architecture, control pulses are generated at more accessible intermediate frequencies around 10 GHz using conventional microwave sources, then upconverted to 300 GHz locally on the chip itself. This approach leverages Josephson mixers and parametric amplifiers, which can provide the necessary $\times 30$ frequency multiplication while maintaining quantum-limited noise performance. The on-chip conversion reduces requirements on room-temperature control electronics, potentially lowering the control stack cost below the €500k estimate for direct 300 GHz generation. However, this approach introduces additional on-chip complexity and requires careful management of spurious mixing products and intermodulation distortion.

2.15.4 Challenge 4: Thermal Management at 4 K

Operating at 4 K rather than the conventional 20 mK introduces requirements for temperature stability below 1 mK, thermal isolation of the qubit chip, and maintenance of a low-noise environment. However, this challenge represents the lowest severity risk, as the engineering principles are well-understood and the technology is mature, resulting in less than 10% probability of failure.

Modern pulse-tube coolers, such as the Cryomech PT420, provide 1 W of cooling power at 4 K with temperature stability better than 0.5 mK. Vibration isolation is achieved through passive damping systems that prevent mechanical noise from coupling to the qubit. Multi-stage thermal filtering at 40 K, 10 K, and 4 K, combined with copper braid thermalization and Eccosorb microwave filters, ensures both thermal and electromagnetic

isolation of the qubit environment. The maturity and reliability of these technologies result in a success probability exceeding 90%.

2.15.5 Comparison Summary

Table 2.20 compares baseline and innovation systems.

Table 2.20: Comparison of baseline and innovation systems

Parameter	Baseline	Innovation	Change
Frequency	5.8 GHz	300 GHz	$52\times$
Temperature	20 mK	4 K	$200\times$
Quantum ratio	13.9	3.6	$0.26\times$
Thermal population	0.000014%	2.81%	$2 \times 10^5\times$
T_1	$80 \mu s$	$20 \mu s$	$0.25\times$
T_2	$60 \mu s$	$14.7 \mu s$	$0.25\times$
F_{single}	99.93%	99.85%	-0.08%
F_{CNOT}	98.7%	97.9%	-0.8%
Cryostat cost	€1.5M	€75k	$0.05\times$
Cooldown time	5 days	4 hours	$0.033\times$

The innovation accepts modest ($< 1\%$) fidelity degradation in exchange for transformative improvements in cost, cooldown time, power, and scalability.




Figure 2.21: **Technical Challenges and Mitigation Strategies.** (a) Loss mechanism scaling comparison between baseline (5.8 GHz, green) and innovation (300 GHz, various colors). Dielectric loss (red) increases $52\times$. Radiative loss (dark red) increases catastrophically by $2675\times$ if unshielded. Quasiparticle loss (green) remains constant—key advantage. Log scale emphasizes dramatic scaling differences. The inset box highlights the critical challenge and the mitigation approach. (b) Substrate comparison at 300 GHz. Standard silicon (red) gives unacceptably short $T_1 \sim 10 \mu\text{s}$. High-resistivity silicon (orange) and sapphire (green) provide $T_1 > 100\mu\text{s}$. Target (dark green) shows achievable performance with optimized materials and design. The red dashed line indicates a $30 \mu\text{s}$ minimum requirement. The green shaded region shows the acceptable range. The inset box recommends sapphire or high- ρ Si. (c) Improvement pathways to reach 98% CNOT threshold. Current system (orange, 97.9%) falls 0.1% short. Individual improvements (green shades): better substrates (+0.3%), lower temperature to 3.5 K (+0.3%), and optimized control pulses (+0.05%). Combined approach (dark green) achieves 99.1%, comfortably exceeding threshold (red dashed line). The green shaded region indicates above-threshold performance. Each bar represents the fidelity value and its improvement over the baseline.

Part D

Conclusions and Recommendations

2.16 Summary of Findings

We have performed a comprehensive analysis of asymmetric SQUID qubits operating at 300 GHz frequency and 4 K temperature. The key findings are:

Technical Feasibility: VIABLE

1. **Quantum regime maintained** at 4 K with $\hbar\omega/k_B T = 3.6$
2. **Manageable thermal population** of 2.81% at 4 K
3. **Adequate coherence times** projected: $T_1 = 20 \mu s$, $T_2 = 14.7 \mu s$
4. **High single-qubit fidelity**: 99.85%
5. **Near-threshold two-qubit fidelity**: 97.9% (0.1% from 98% target)

Baseline Validation: CONFIRMED

- 5.8 GHz / 20 mK system achieves 98.7% CNOT fidelity
- Validates asymmetric SQUID architecture and design methodology
- Provides confidence in theoretical models
- De-risks innovation phase

Technical Challenges: ADDRESSABLE

Four major challenges identified, all with credible mitigation strategies:

1. **Material losses** (dielectric, radiative): Ultra-low-loss substrates + 3D cavity
2. **Junction fabrication**: Advanced lithography + high current density
3. **300 GHz control**: Leverage mm-wave astronomy technology
4. **Thermal management**: Modern pulse-tube coolers

Overall technical success probability: 10–15%

Economic Impact: TRANSFORMATIVE

- **10× cost reduction:** €75k vs €750k–€1.5M
- **24–100× faster development:** Hours vs days cooldown
- **10× lower power:** 0.5 kW vs 5 kW
- **True scalability:** 10,000+ qubits per cooler vs 1000
- **Market expansion:** €300M–€5B value creation potential

The proposed 300 GHz / 4 K asymmetric SQUID qubit represents a **high-risk, high-reward** innovation at the frontier of quantum engineering. While technical challenges are significant, they are addressable with existing technology and clear mitigation strategies.

The validated baseline performance (98.7% CNOT at conventional conditions) provides crucial confidence that the asymmetric SQUID architecture and our theoretical framework are sound. The innovation "merely" requires maintaining most of this performance while operating at elevated temperature—a difficult but not impossible goal.

The economic and scalability benefits are transformative: 10× cost reduction, 100× faster development cycles, and path to 10,000+ qubit processors. These advances would democratize quantum computing, expanding the addressable market by orders of magnitude.

The incremental four-phase approach, with go/no-go decision points, manages risk appropriately. Even partial success (e.g., single qubits at 4 K) would represent a significant scientific breakthrough.

We therefore recommend **proceeding with Phase 1**, targeting demonstration of 300 GHz qubits at baseline conditions within 12–18 months.

2.17 Wrap-up and Final Considerations

Part A: Baseline (Proven)

- Asymmetric SQUID qubits at 5–6 GHz achieve 98.7% CNOT fidelity at 20 mK
- Thermal effects negligible ($< 10^{-5}\%$)
- Competitive with state-of-the-art
- **Demonstrates the concept works**

Part B: Innovation (Novel)

- 300 GHz operation enables 4–10 K temperatures
- Thermal populations manageable (3–10%)
- Projected fidelities above error correction thresholds
- **Pathway to transformative breakthrough**

State-of-the-art superconducting qubits use Al/AlO_x/Al junctions operated at 10–20 mK in dilution refrigerators [2], achieving coherence times exceeding 300 μ s. In contrast, our approach employs epitaxial NbN/AlN/NbN junctions on sapphire [78, 79], which have demonstrated $T_1 \approx 16 \mu$ s and $T_c \approx 16$ K at conventional operating temperatures, and we project these to enable coherent operation at 4–10 K in compact cryocoolers. This opens the path to direct co-integration with cryo-CMOS electronics [119, 120, 121], drastically reducing cryogenic infrastructure cost (70–80% savings), cooldown time, and system footprint.

What is new vs. prior art. We demonstrate Kelvin-scale coherent qubits with engineered junction asymmetry for tunability and flux-noise suppression. Operation at ~ 10 K maintains transmon anharmonicity ($|\alpha|/2\pi \sim 0.1\text{--}0.3$ GHz) and dispersive readout ($\chi/\kappa \gg 1$), matching circuit-QED requirements.

Why now. Advances in high- T_c nitride epitaxy, robust microwave control hardware, and commercial cryocoolers enable practical operation above 4 K. The possibility of seamless cryo-CMOS integration motivates a transition beyond dilution refrigerators.

Competitive advantage.

1. *Cost/infrastructure:* removal of dilution refrigerators cuts capex/opex by 70–80%.
2. *Integration:* short control chains reduce attenuation, latency, and wiring complexity.
3. *Throughput:* cooldown in hours (vs. days) accelerates fab–test iterations.
4. *Scalability:* 2D layouts remain compatible with surface-code error correction [12, 122].

All-nitride epitaxial junctions and tantalum-based films [73, 78] extend coherence to the millisecond range. NbN technology enables operation above liquid-helium temperature, easing cryogenic integration while preserving strong nonlinearity. Coupling with Rydberg-based receivers may enable hybrid quantum sensing–computing architectures. We presented a design framework for a 4-K asymmetric SQUID transmon using NbN/AlN/NbN junctions, discussed losses and decoherence, and outlined EM and readout strategies. The approach bridges superconducting hardware and quantum sensing, supporting the integration of these technologies within scalable hybrid quantum systems.

2.17.1 Technological prerequisites for high-fidelity 4 K operation

The cross-resonance protocol of §2.11 reduces the two-qubit gate fidelity question to a clear set of materials and engineering requirements. We identify here, with quantitative targets, the six prerequisites for reaching $F_{\text{avg}} \geq 99\%$ on the innovation HEATS-Q operating point ($\omega_q \sim 300$ GHz, $T = 4$ K).

P1: Thermal occupation. The qubit frequency must satisfy $\omega_q \gtrsim 250$ GHz so that $\bar{n}_{\text{th}}(\omega_q, 4 \text{ K}) < 0.05$. At $\omega_q = 300$ GHz the thermal population is $\bar{n}_{\text{th}} = 2.8\%$, marginal but acceptable provided that active reset is employed at the start of every circuit [81].

P2: Dielectric loss tangent. At $\omega_q = 300$ GHz, achieving $T_1 \geq 25 \mu$ s (the threshold for $F_{\text{avg}} \geq 99\%$, see Fig. 2.22) requires $\tan \delta \leq 1/(\omega_q T_1) \simeq 2 \times 10^{-8}$. This figure is at the optimistic end of state-of-the-art microwave dielectric measurements ([113, 117, 118], all reporting $\tan \delta \sim 10^{-8}$ for ultra-pure single-crystal silicon and sapphire), but is not unreasonable given the sub-mK $\tan \delta$ values achieved in the bulk-niobium 3D cavities




Figure 2.22: Achievable cross-resonance gate fidelity at $T = 4$ K as a function of qubit frequency ω_q and coherence time T_1 , with Δ_q independently optimized at each point. The bold solid contour at $F_{\text{avg}} = 99\%$ defines the boundary of the operational sweet spot. The red cross marks the original-manuscript point ($T_1 = 5 \mu\text{s}$); the blue star marks the redesigned point ($T_1 = 50 \mu\text{s}$) discussed in this section. The gate time, optimized over Δ_q , is approximately constant (~ 200 ns) over the parameter range shown (right panel).

of Ref. [114]. A direct experimental verification of $\tan \delta$ at sub-THz frequencies in 4 K cavities is, to our knowledge, an open question and is identified here as a key prerequisite.

P3: Junction superconducting gap above $\hbar\omega_q$. Quasi-particle dynamics, the dominant non-radiative decoherence channel for high-frequency qubits at $T \gtrsim 1$ K, is exponentially suppressed when $\hbar\omega_q < 2\Delta_{\text{gap}}$. The NbN/AlN/NbN junctions of Refs. [79, 123] have $2\Delta_{\text{gap}}/h \approx 720$ GHz, well above $\omega_q = 300$ GHz. This places the NbN junction technology at the centre of the HEATS-Q design.

P4: Stretched-CR detuning regime. Following §2.11.2, the optimal qubit–qubit detuning is $\Delta_q \approx 600$ MHz (innovation point), corresponding to qubit frequencies $\omega_{q,1}/2\pi = 300.30$ GHz and $\omega_{q,2}/2\pi = 299.70$ GHz. This is a redesign of the original $\Delta_q = 20$ GHz value, motivated by the need to keep the prefactor in Eq. (2.112) large; spectral addressability with single-qubit pulses of duration $\tau \sim 20$ ns is preserved, since Δ_q exceeds the pulse bandwidth ~ 50 MHz by an order of magnitude.

P5: Echo-CR pulse engineering. The echo-CR sequence [112], complemented by Gaussian-Flat-Gaussian shaping with DRAG correction [5, 103], brings the gate fidelity from the bare-Hamiltonian $\sim 80\%$ to the surface-code-compatible $\gtrsim 99\%$ (see Table 2.12). All necessary techniques are routinely employed in current transmon processors at GHz frequencies; their extension to sub-THz drives is a control-electronics challenge addressed in §2.2.6.

P6: Cavity quality factor at 4 K. The Purcell limit [82] on T_1 requires $Q_{\text{cav}} \gtrsim 10^4$ at the operating frequency. For 3D cavities of dimensions consistent with $\omega_r = 320$ GHz, this is within the demonstrated capability of high-purity bulk-Nb fabrication [114] but

has not yet been demonstrated at $\omega \gtrsim 100$ GHz and $T = 4$ K. This represents the second key open experimental question, in addition to P2.

In summary, the high-temperature operation of HEATS-Q is shown to be feasible within a sweet spot whose boundary is defined by quantitatively explicit and independently testable engineering targets (P1–P6). Two prerequisites (P2 and P6) require dedicated experimental validation at sub-THz frequencies, while the others (P1, P3, P4, P5) draw on technology already demonstrated at lower frequencies. The path to a high-fidelity universal gate set at 4 K, while challenging, does not invoke any speculative physics: it is a problem of materials and pulse engineering whose components have, individually, been demonstrated.